

EECE 2322: Fundamentals of Digital Design and Computer Organization

Lecture 15_1: Microarchitecture

Xiaolin Xu

Department of ECE
Northeastern University

Review: Exceptions

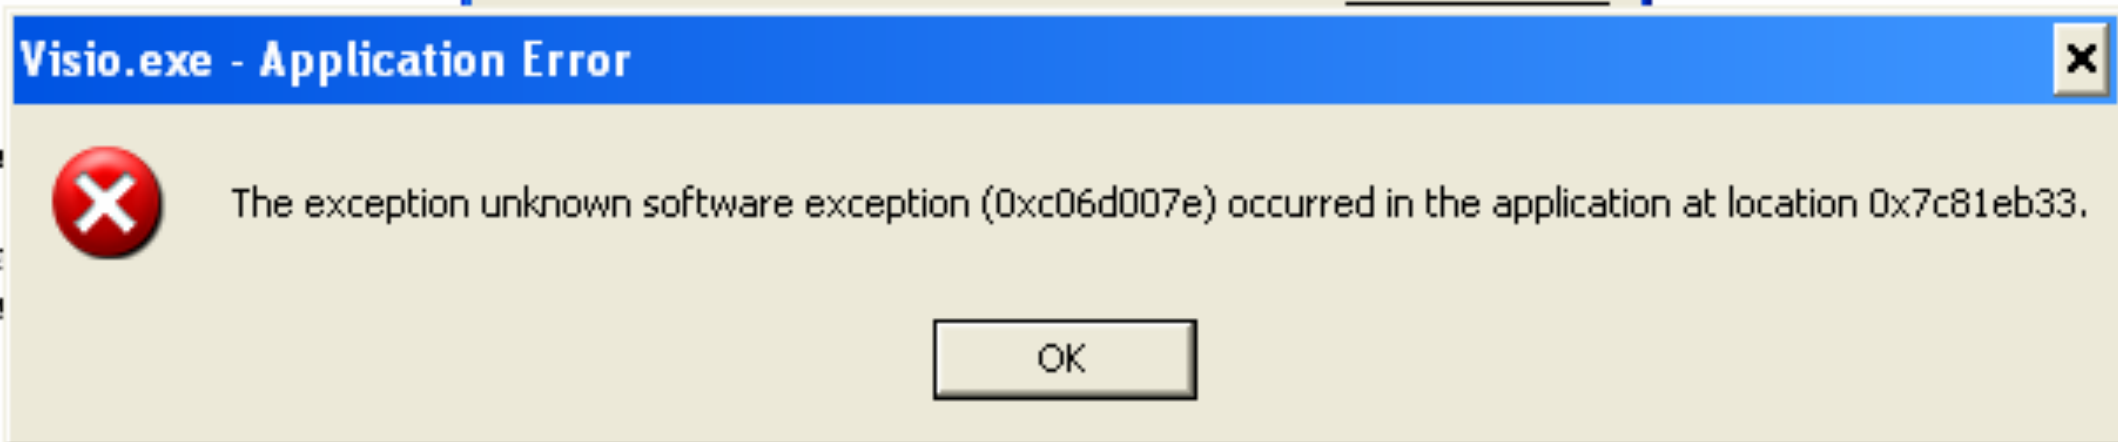
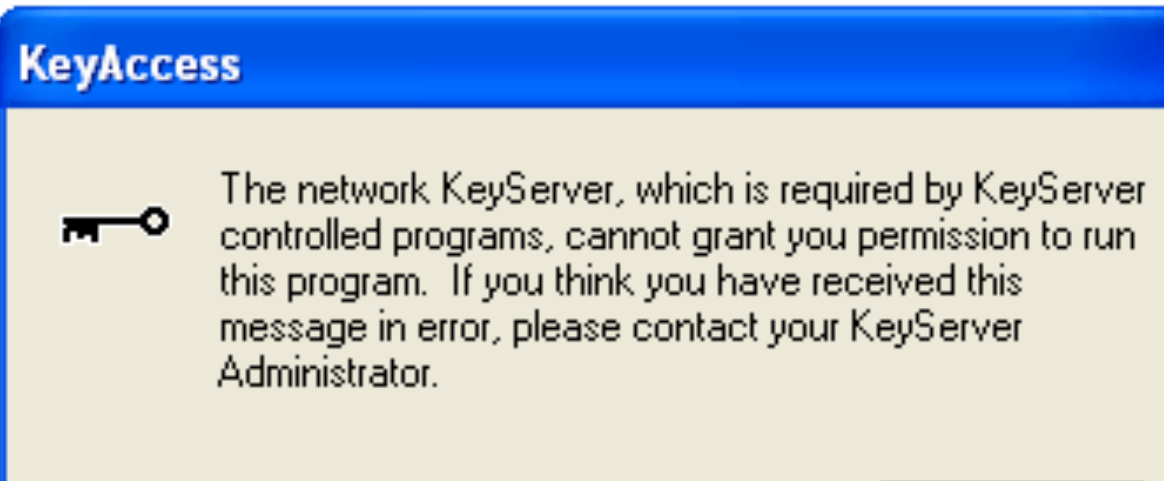
- Unscheduled function call to *exception handler*
- Caused by:
 - Hardware, also called an *interrupt*, e.g. keyboard
 - Software, also called *traps*, e.g. undefined instruction
- When exception occurs, the processor:
 - Records cause of exception (Cause register)
 - Jumps to exception handler (0x80000180)
 - Returns to program (EPC register)

Example Exception

sequential circuits.

Can we design a spiff

Figure 2.11 shows a inputs, A and B, and on box indicates that it is this case, the function is



words, we say the output Y is a function of the two inputs A and B where the function performed is A OR B.

The *implementation* of the combinational circuit is independent of its functionality. Figure 2.1 and Figure 2.2 show two possible implementa-

Exception Registers

- ❖ Not part of register file
 - ❖ Cause
 - ❖ Records cause of exception
 - ❖ Coprocessor 0 register 13
 - ❖ EPC (Exception PC)
 - ❖ Records PC where exception occurred
 - ❖ Coprocessor 0 register 14

- ❖ Move from Coprocessor 0

mfc0

- ❖ `mfc0 $t0, Cause`

010000	00000	\$t0 (8)	Cause (13)	000000000000
31:26	25:21	20:16	15:11	10:0

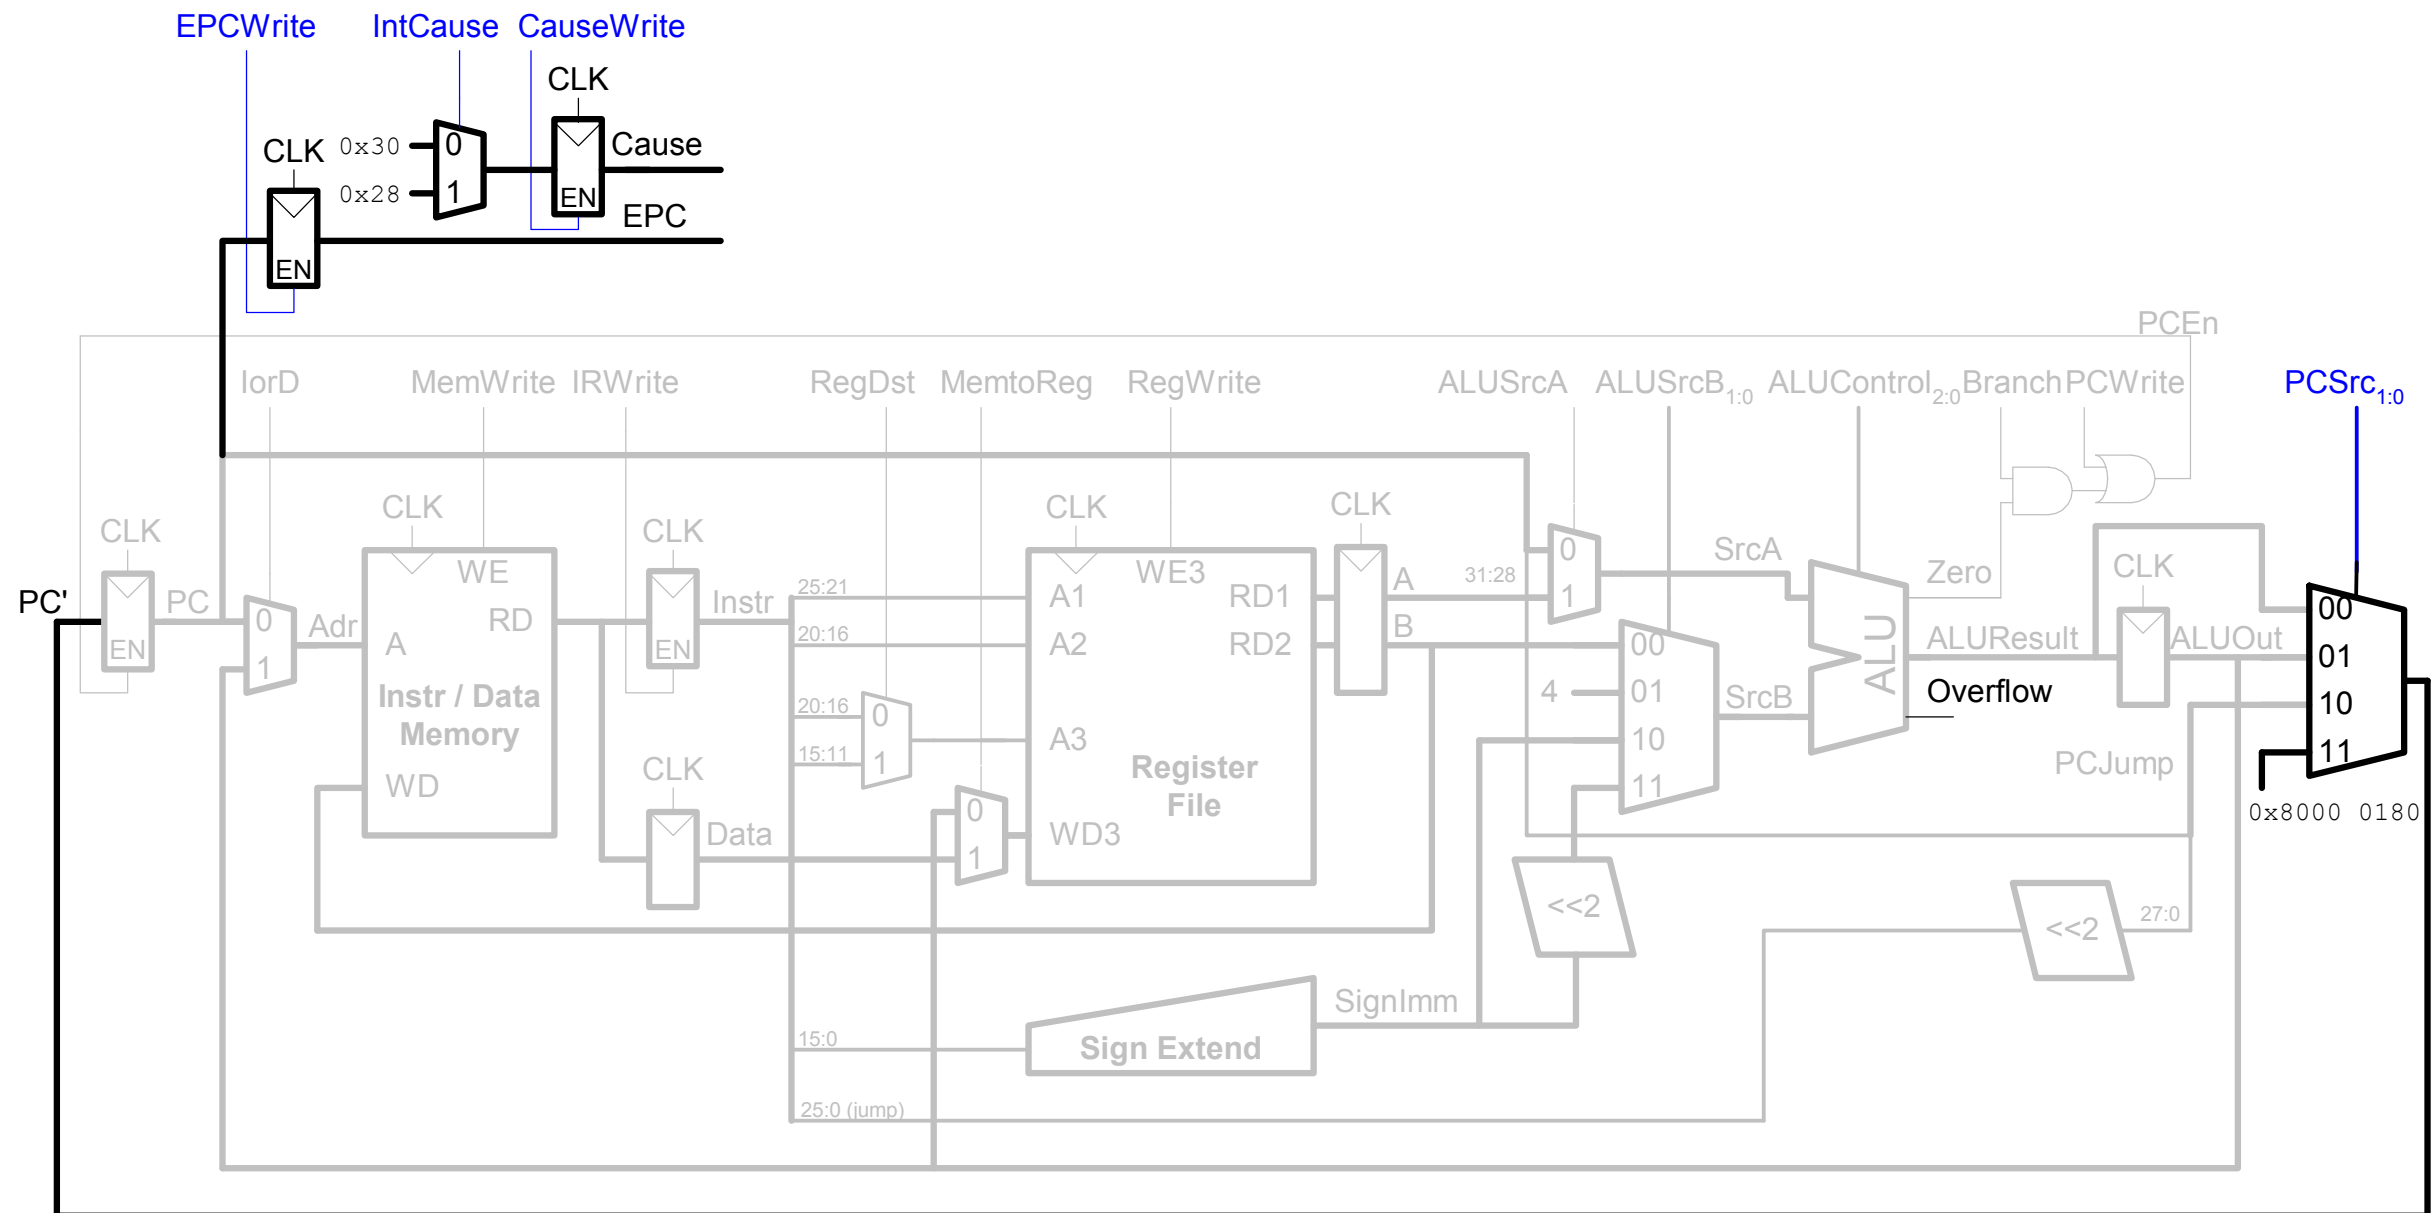
- ❖ Moves contents of Cause into `$t0`

Exception Causes

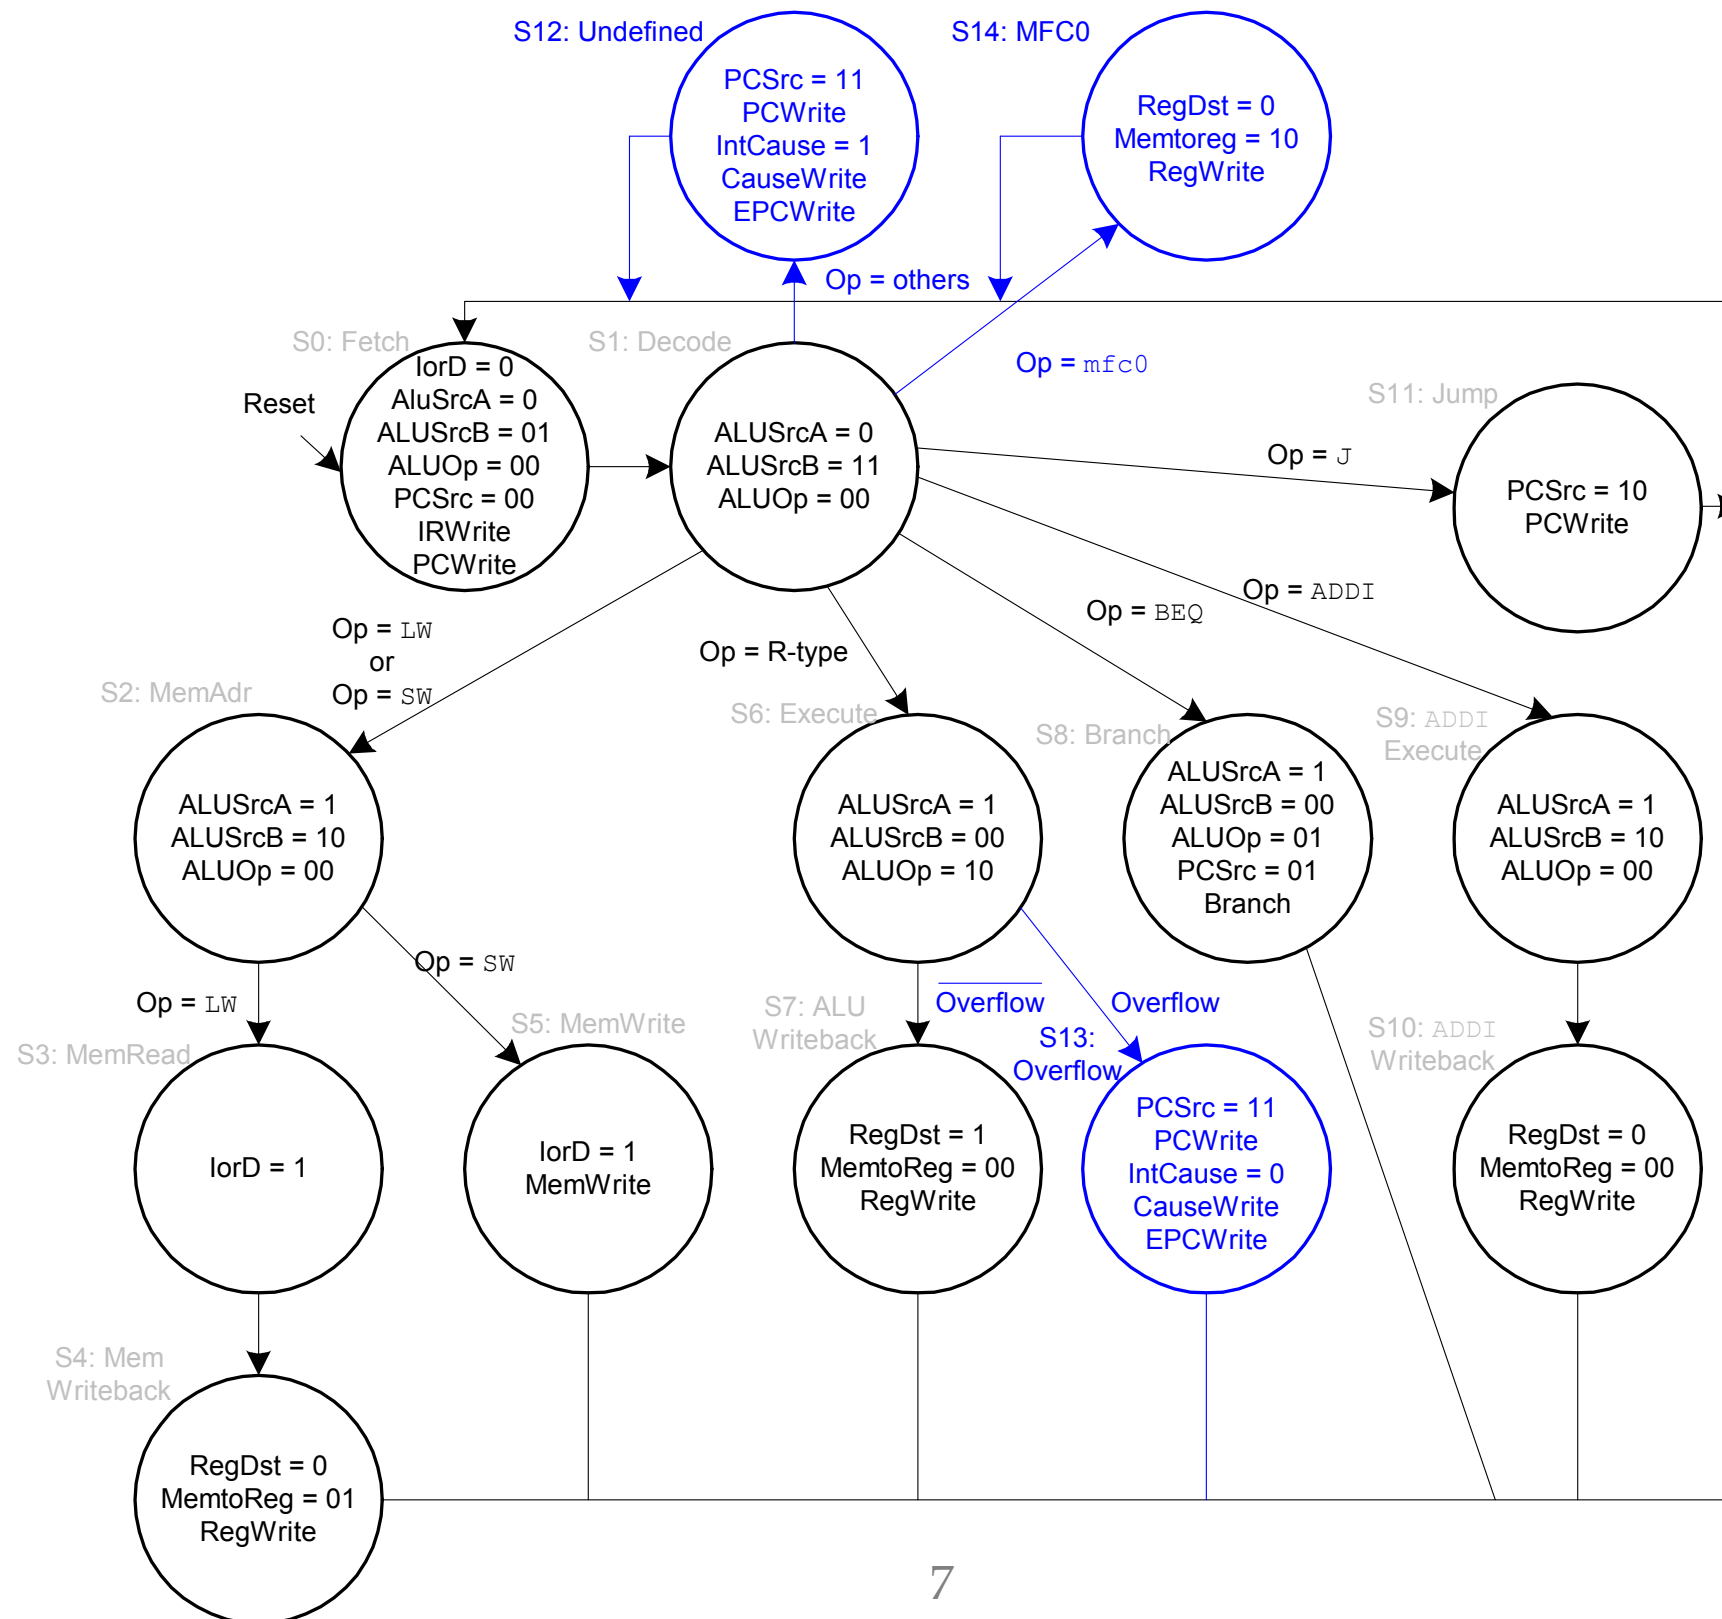
Exception	Cause
Hardware Interrupt	0x00000000
System Call	0x00000020
Breakpoint / Divide by 0	0x00000024
Undefined Instruction	0x00000028
Arithmetic Overflow	0x00000030

Extend multicycle MIPS processor to handle last two types of exceptions

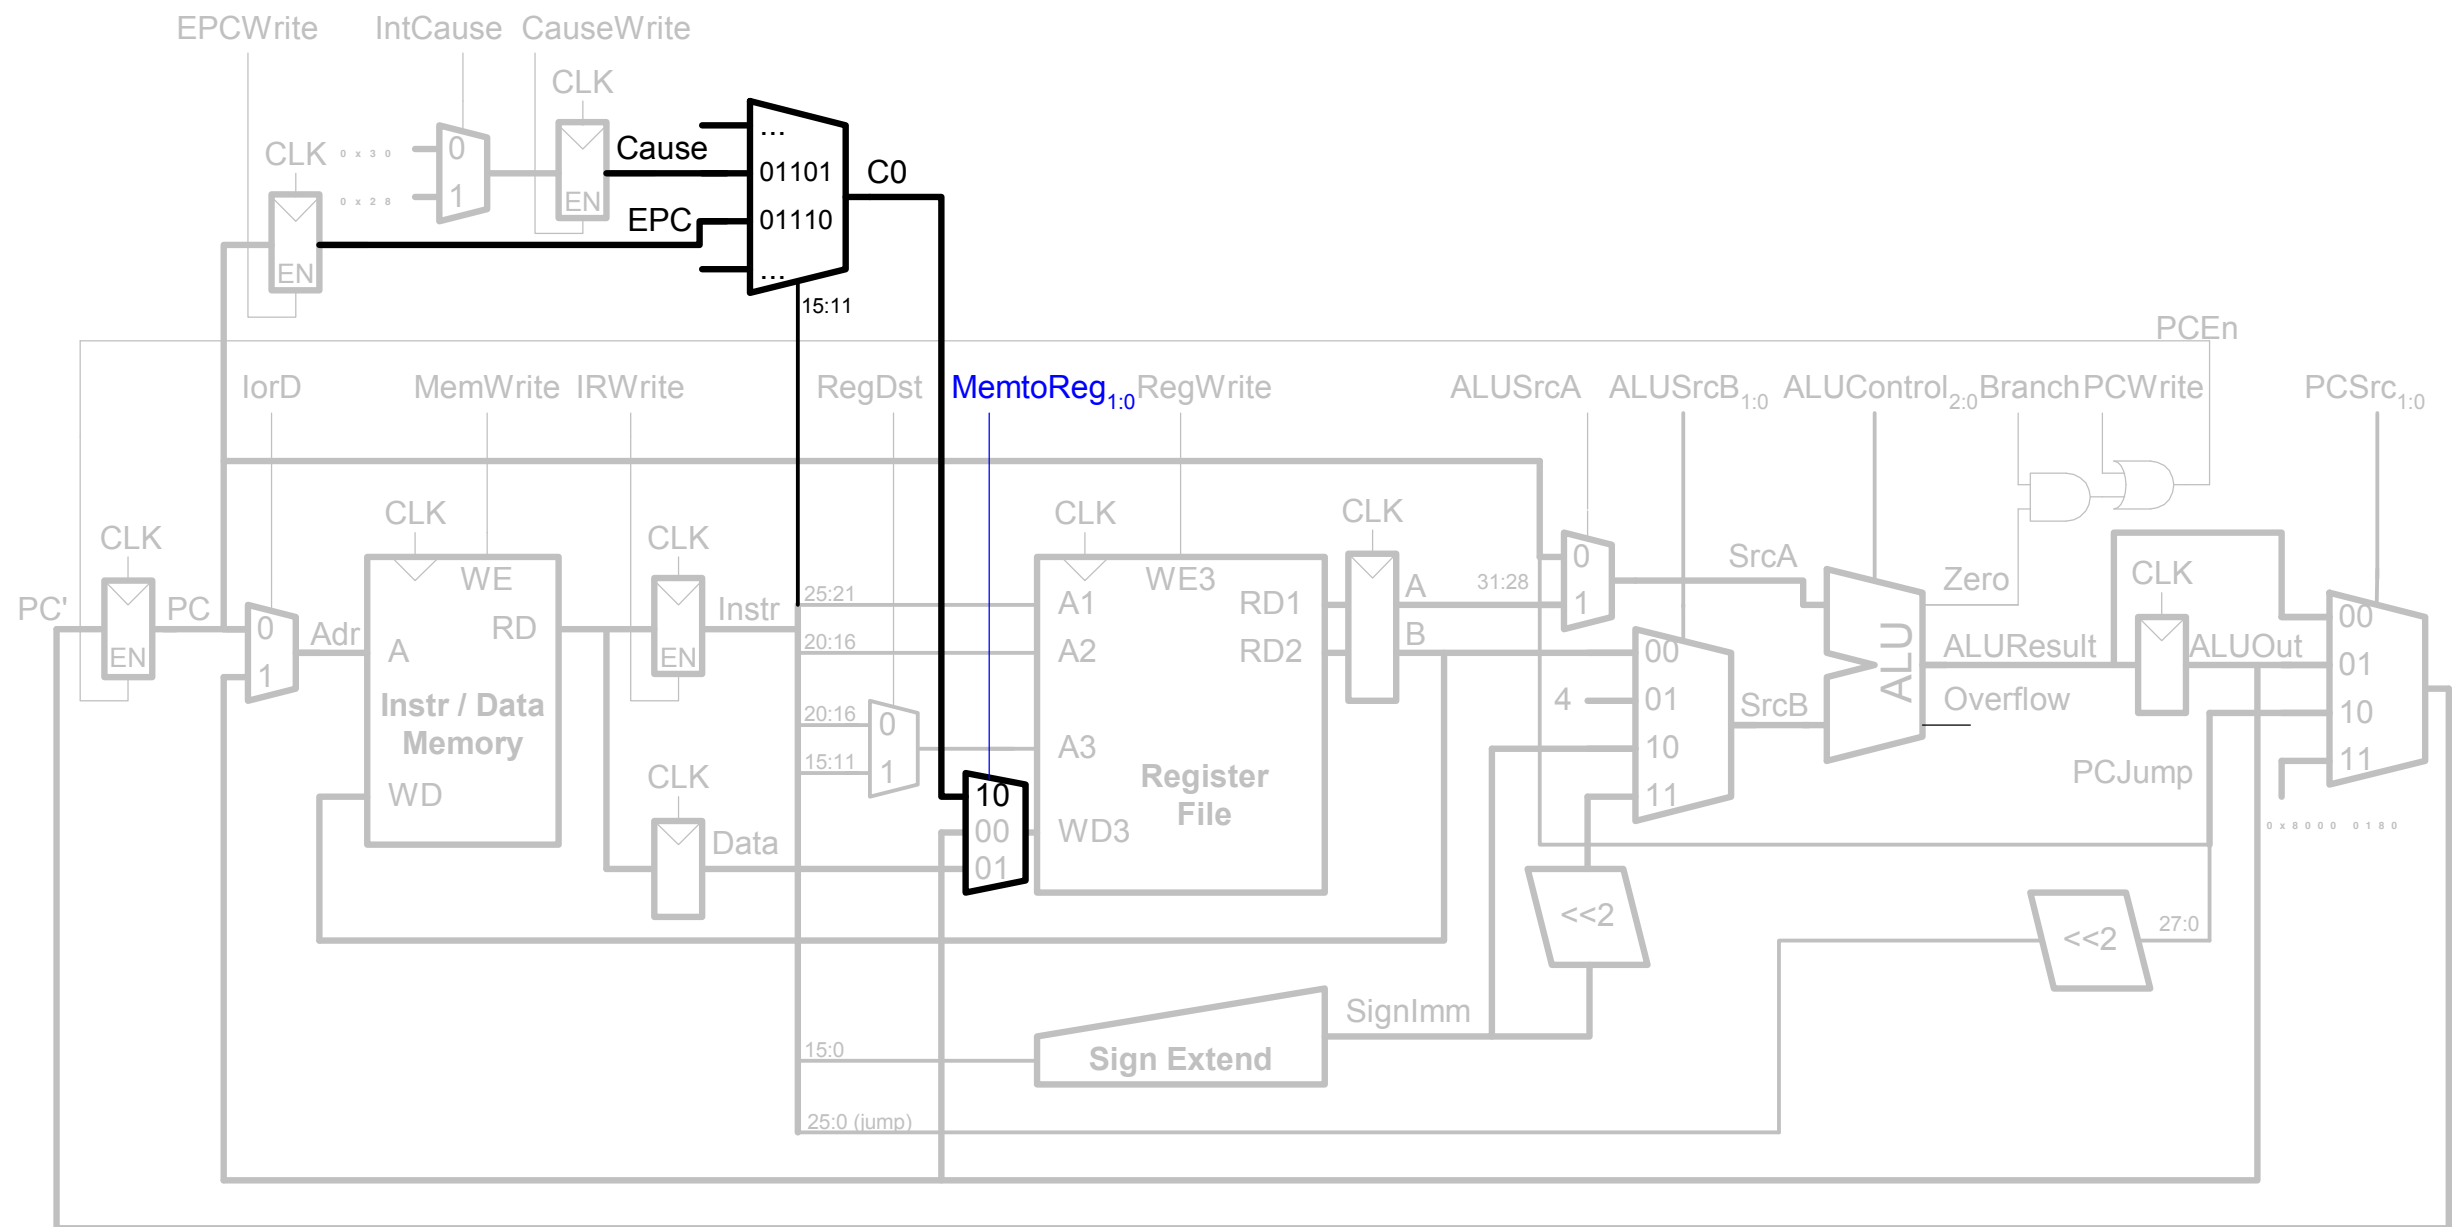
Exception Hardware: EPC & Cause



Control FSM with Exceptions



Exception Hardware: mfc0



Advanced Microarchitecture

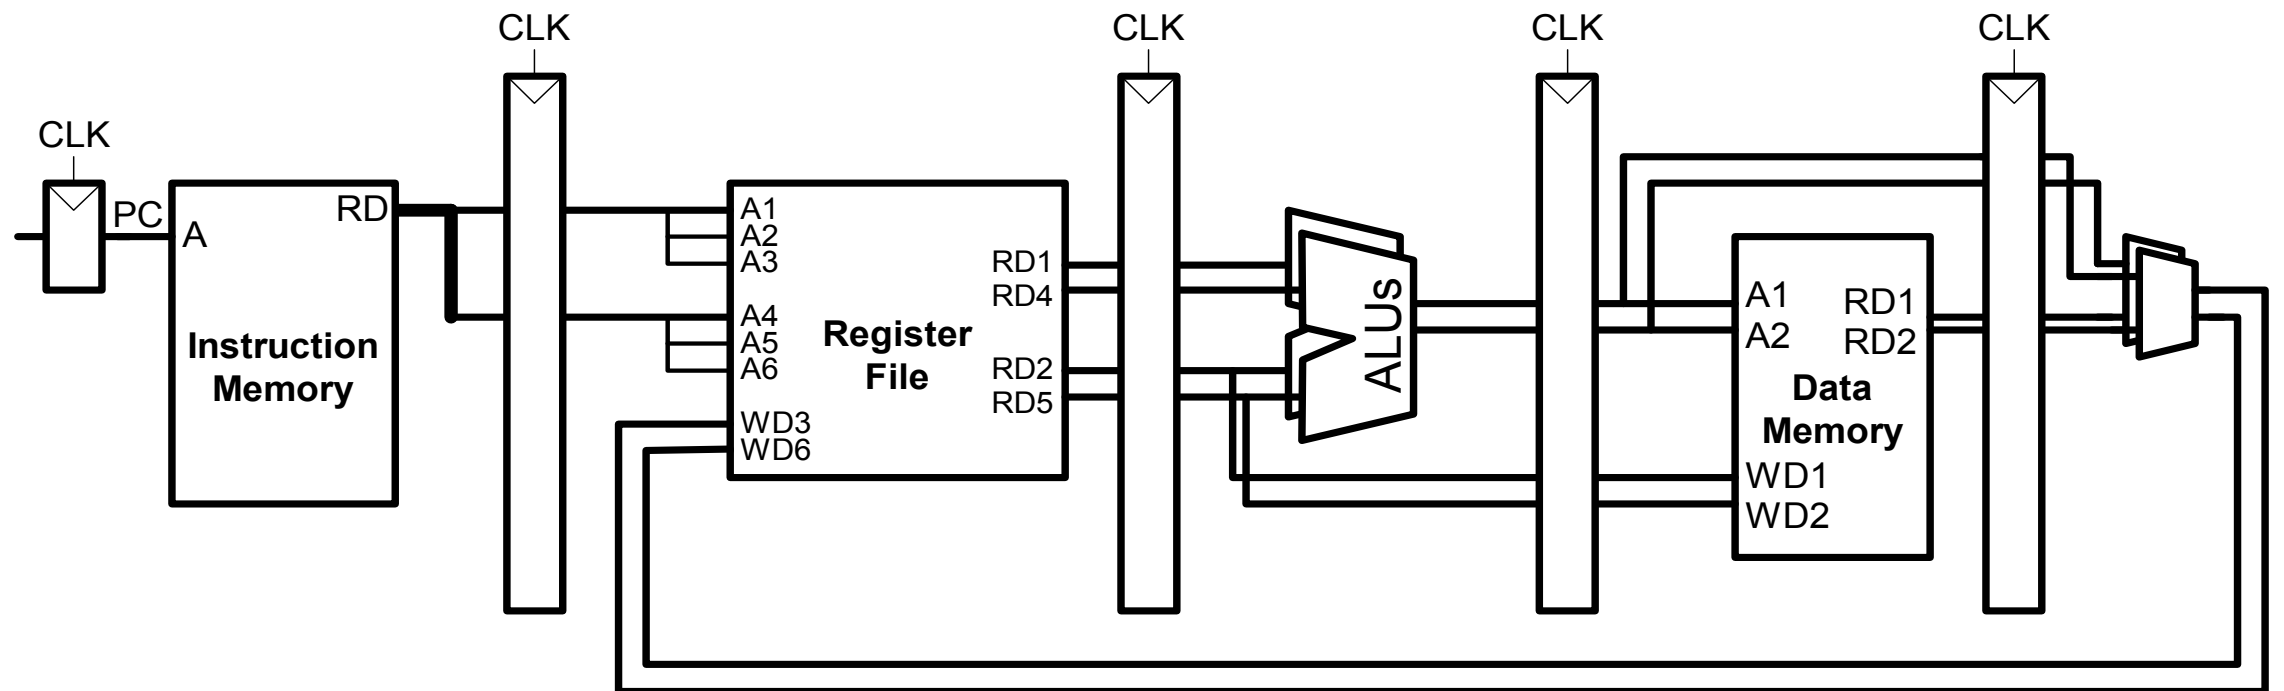
- ❖ Deep Pipelining
- ❖ Branch Prediction
- ❖ Superscalar Processors
- ❖ Out of Order Processors
- ❖ Register Renaming
- ❖ SIMD
- ❖ Multithreading
- ❖ Multiprocessors

Deep Pipelining

- ❖ 10-20 stages typical
- ❖ Number of stages limited by:
 - ❖ Pipeline hazards
 - ❖ Sequencing overhead
 - ❖ Power
 - ❖ Cost

Superscalar

- ❖ Multiple copies of datapath execute multiple instructions at once
- ❖ Dependencies make it tricky to issue multiple instructions at once

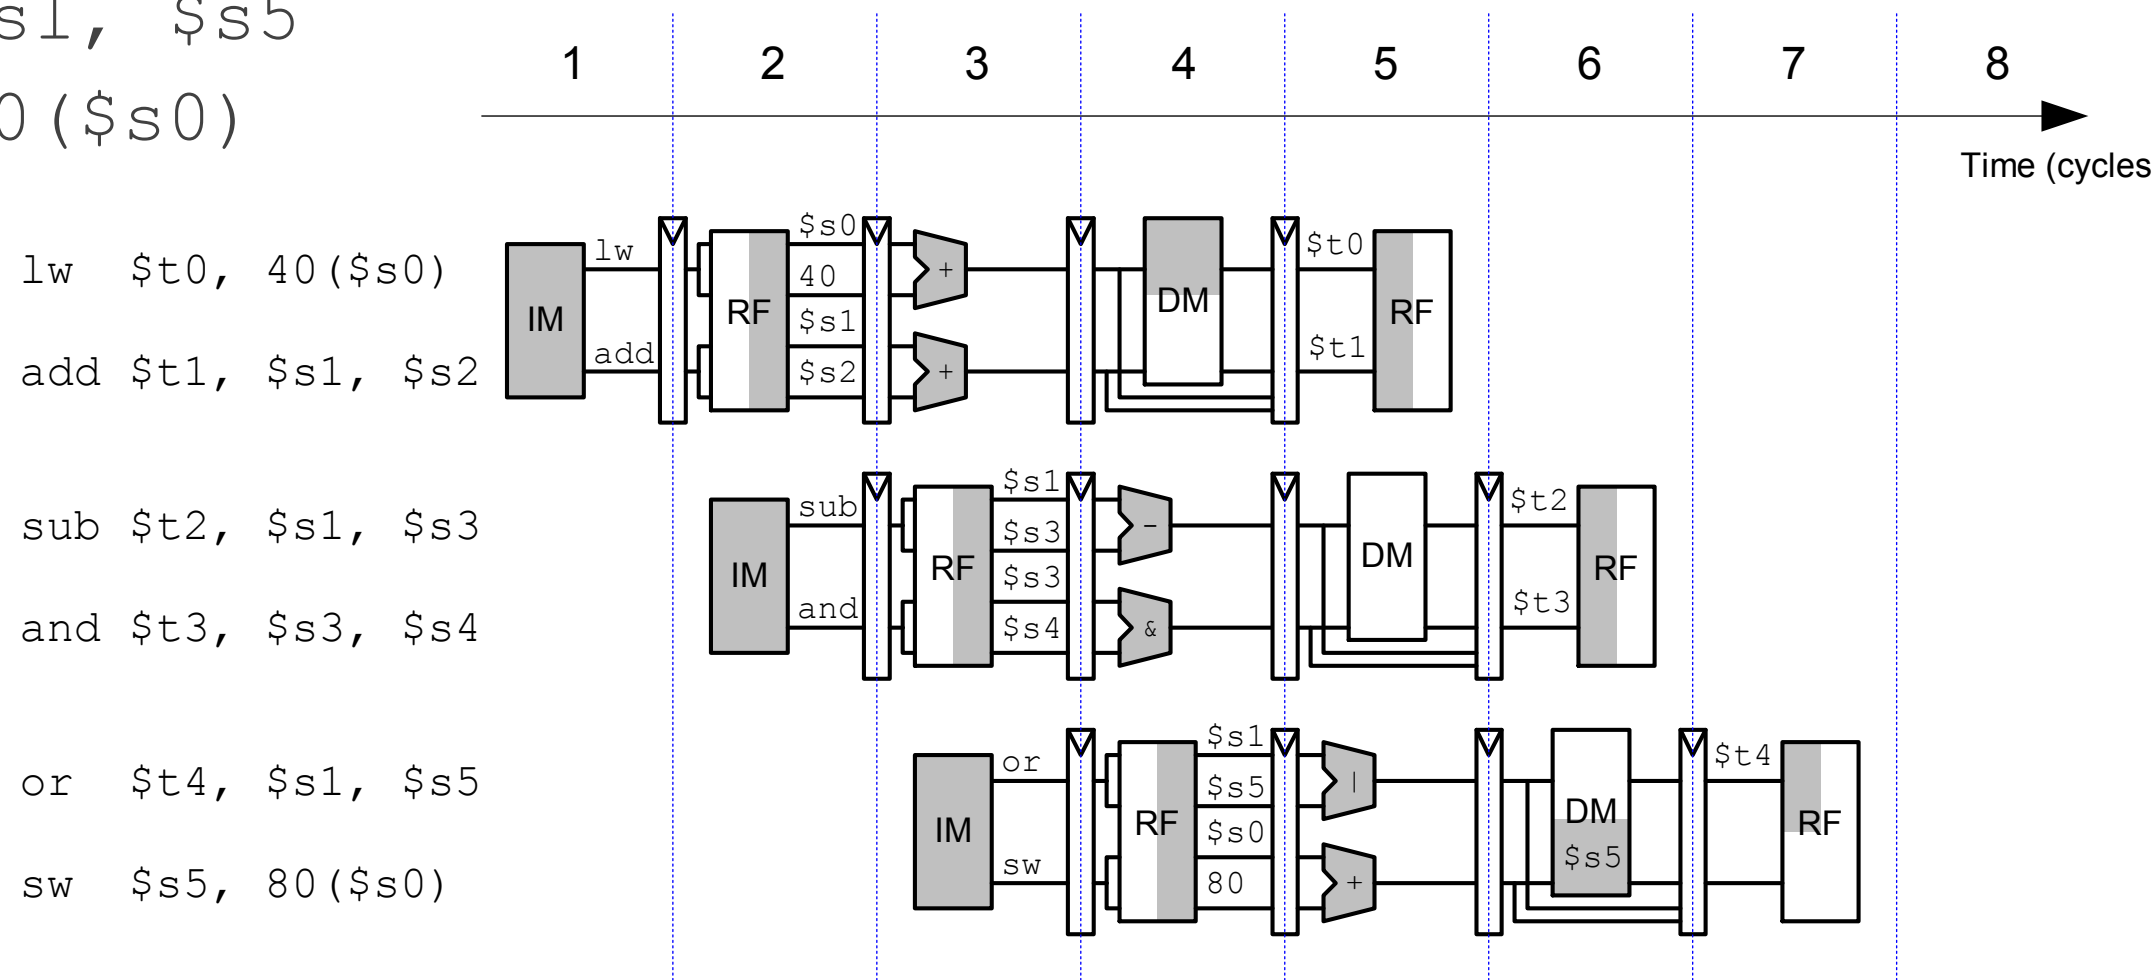


Superscalar Example

```
lw    $t0, 40($s0)
add   $t1, $s1, $s2
sub   $t2, $s1, $s3
and   $t3, $s3, $s4
or    $t4, $s1, $s5
sw    $s5, 80($s0)
```

Ideal IPC (Instruction Per Cycle): 2

Practical IPC = ?



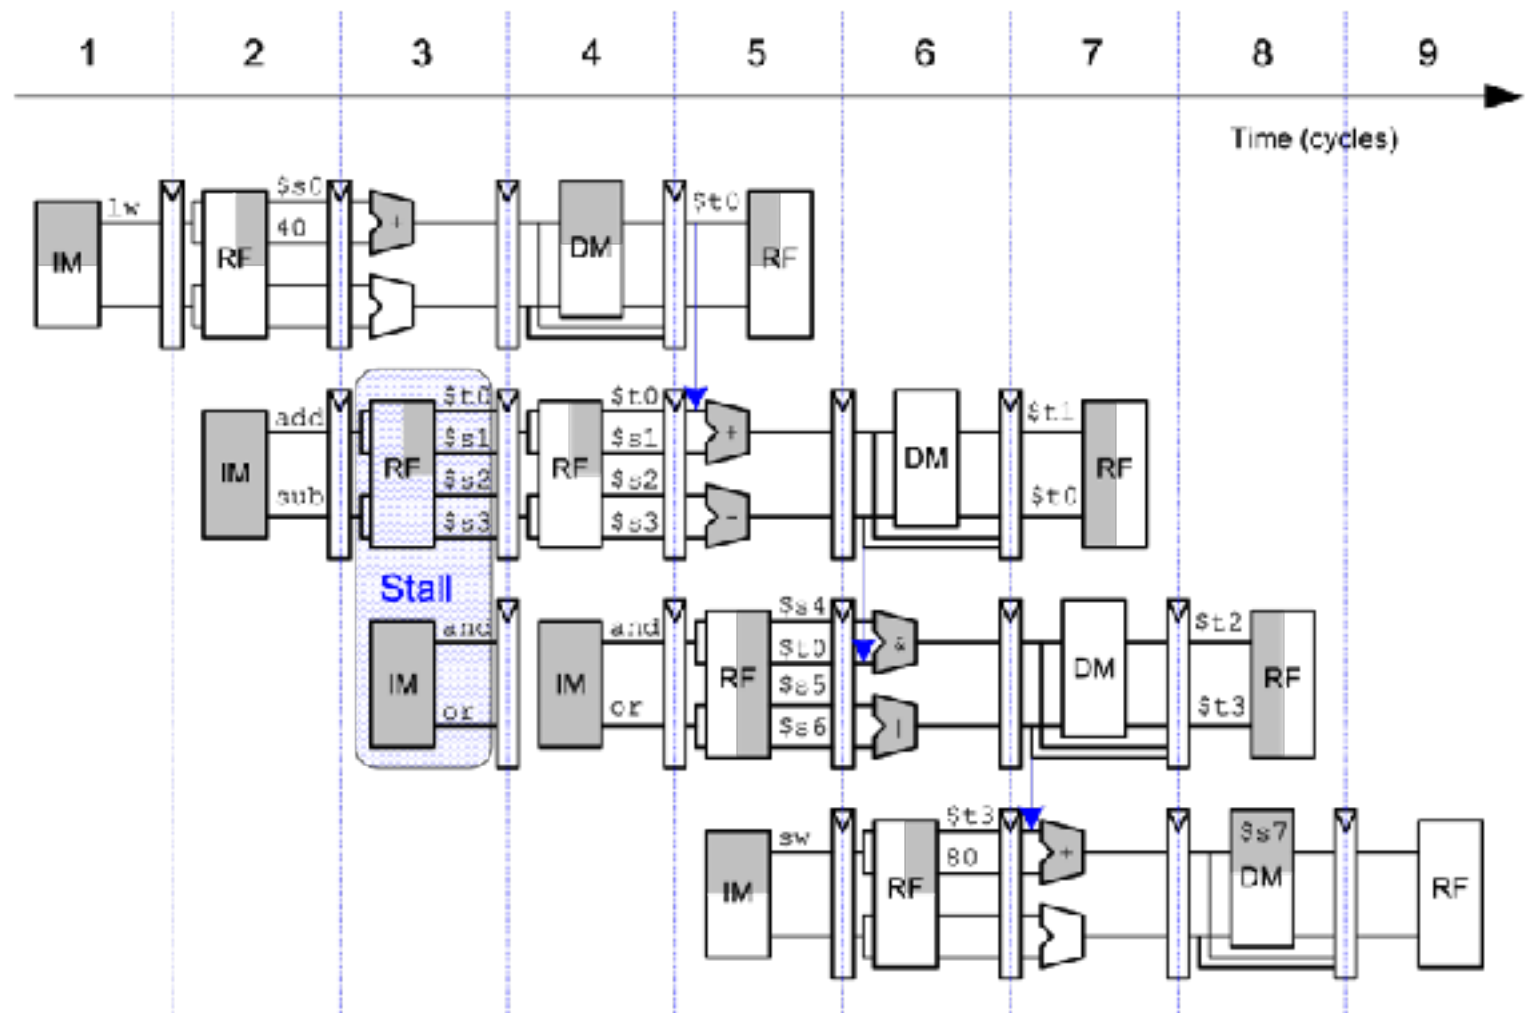
Superscalar with Dependencies

```
lw    $t0, 40($s0)
add   $t1, $t0, $s1
sub   $t0, $s2, $s3
and   $t2, $s4, $t0
or    $t3, $s5, $s6
sw    $s7, 80($t3)
```

Ideal IPC: 2

Actual IPC: $6/5 = 1.17$

lw \$t0, 40(\$s0)
add \$t1, \$t0, \$s1
sub \$t0, \$s2, \$s3
and \$t2, \$s4, \$t0
or \$t3, \$s5, \$s6
sw \$s7, 80(\$t3)



Out of Order Processor

- ❖ Looks ahead across multiple instructions
- ❖ Issue as many instructions as possible at once
- ❖ Issue instructions out of order (**as long as no dependencies**)
- ❖ **Dependencies:**
 - ❖ **RAW** (read after write): one instruction writes, later instruction reads a register, (e.g., lw instruction)
 - ❖ **WAR** (write after read): one instruction reads, later instruction writes a register
 - ❖ **WAW** (write after write): one instruction writes, later instruction writes a register

Out of Order Processor

- ❖ **Instruction level parallelism (ILP):** number of instruction that can be issued simultaneously (average < 3)
- ❖ **Scoreboard:** table that keeps track of:
 - ❖ Instructions waiting to issue
 - ❖ Available functional units
 - ❖ Dependencies

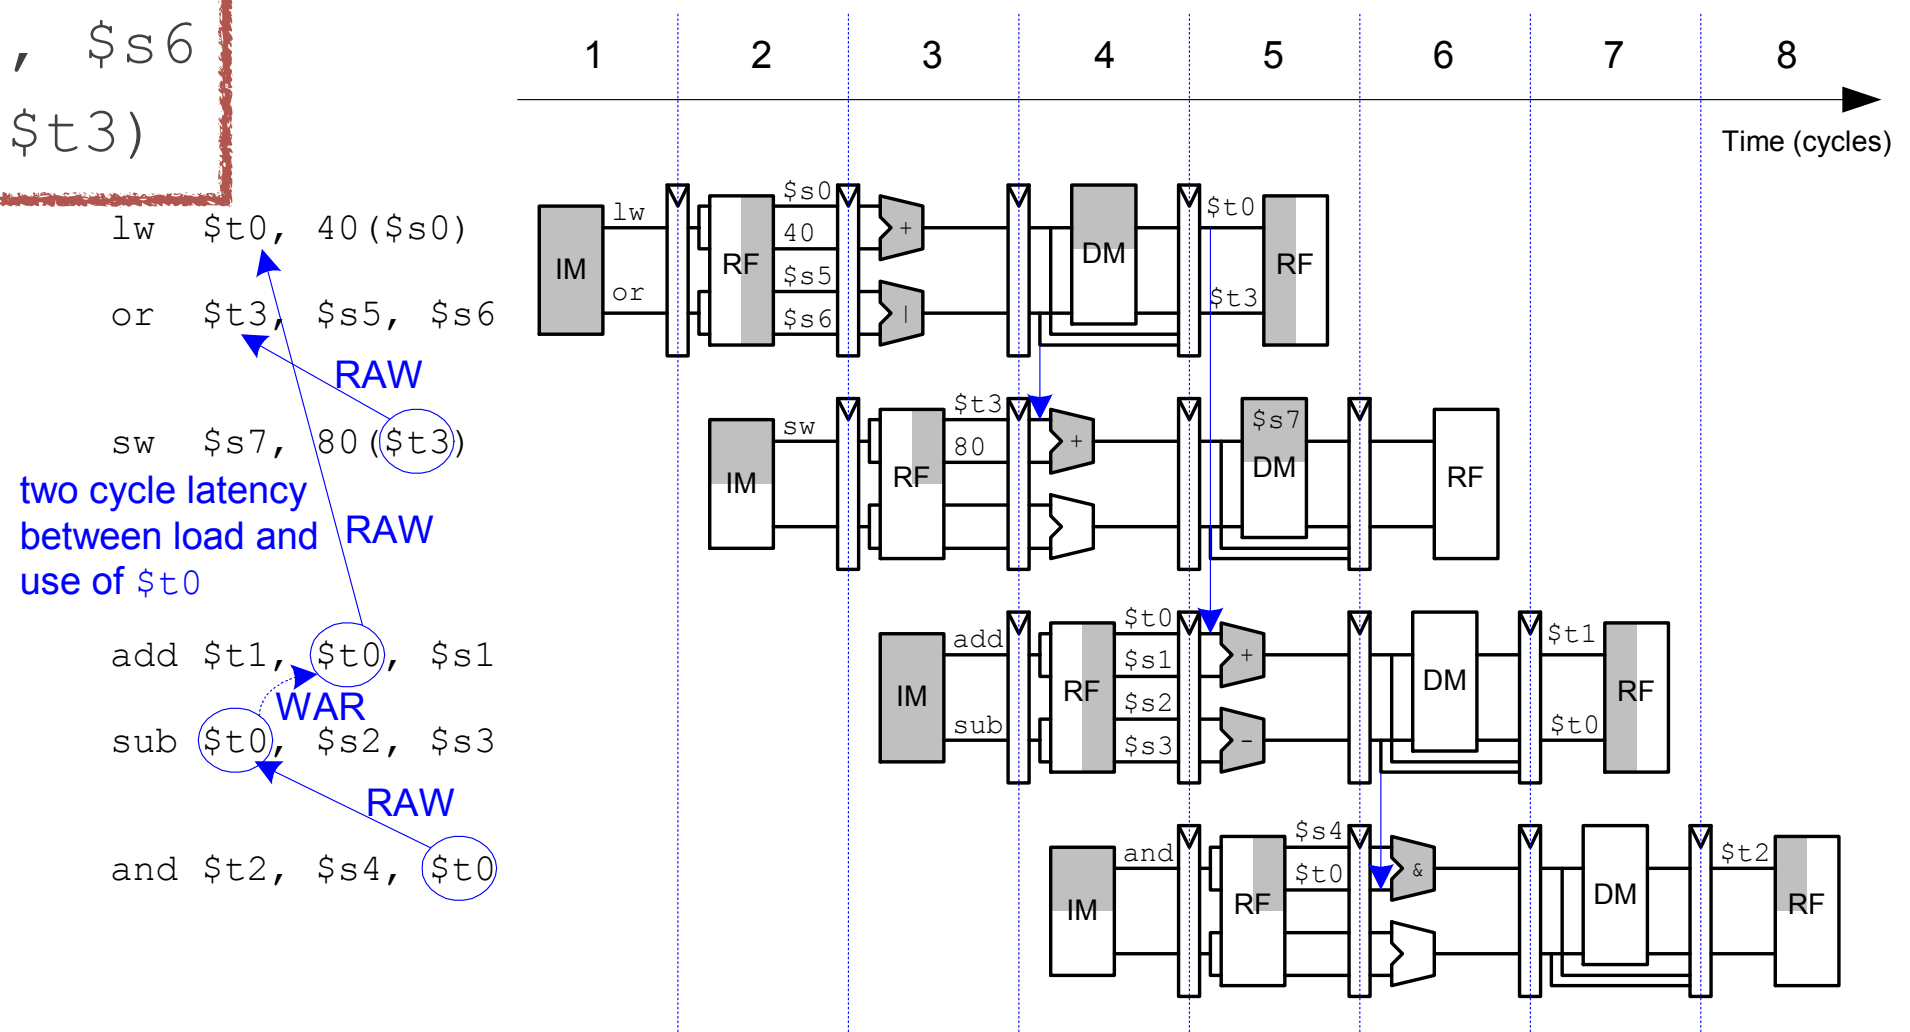
Out of Order Processor Example

original code

```
lw    $t0, 40($s0)
add   $t1, $t0, $s1
sub   $t0, $s2, $s3
and   $t2, $s4, $t0
or    $t3, $s5, $s6
sw    $s7, 80($t3)
```

Ideal IPC: 2

Actual IPC: $6/4 = 1.5$



Out of Order Processor Example

original code

```
lw    $t0, 40($s0)
add   $t1, $t0, $s1
sub   $t0, $s2, $s3
and   $t2, $s4, $t0
or    $t3, $s5, $s6
sw    $s7, 80($t3)
```

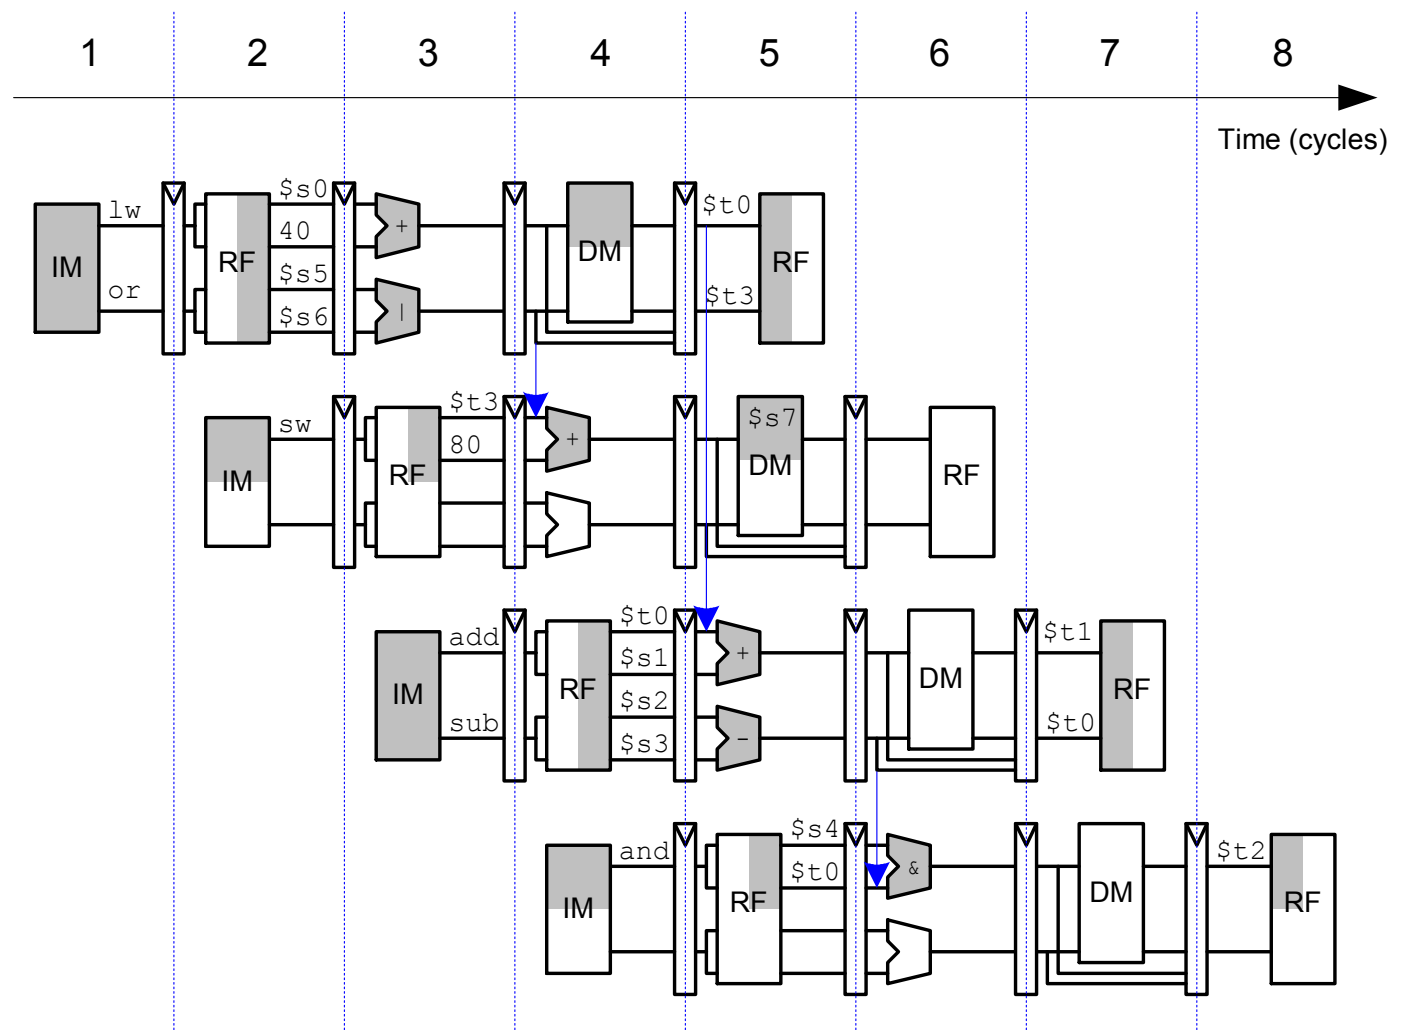
Ideal IPC: 2

Actual IPC: $6/4 = 1.5$

lw \$t0, 40(\$s0)
or \$t3, \$s5, \$s6
sw \$s7, 80(\$t3)
two cycle latency
between load and
use of \$t0

easier to solve

```
add $t1, $t0, $s1
sub $t0, $s2, $s3
and $t2, $s4, $t0
```



Register Renaming

- ❖ MIPS processor might add 20 renaming registers, called \$r0-\$r19
- ❖ The programmer cannot use these registers **directly**, because they are not part of the architecture
- ❖ However, the processor is free to use them to eliminate hazards

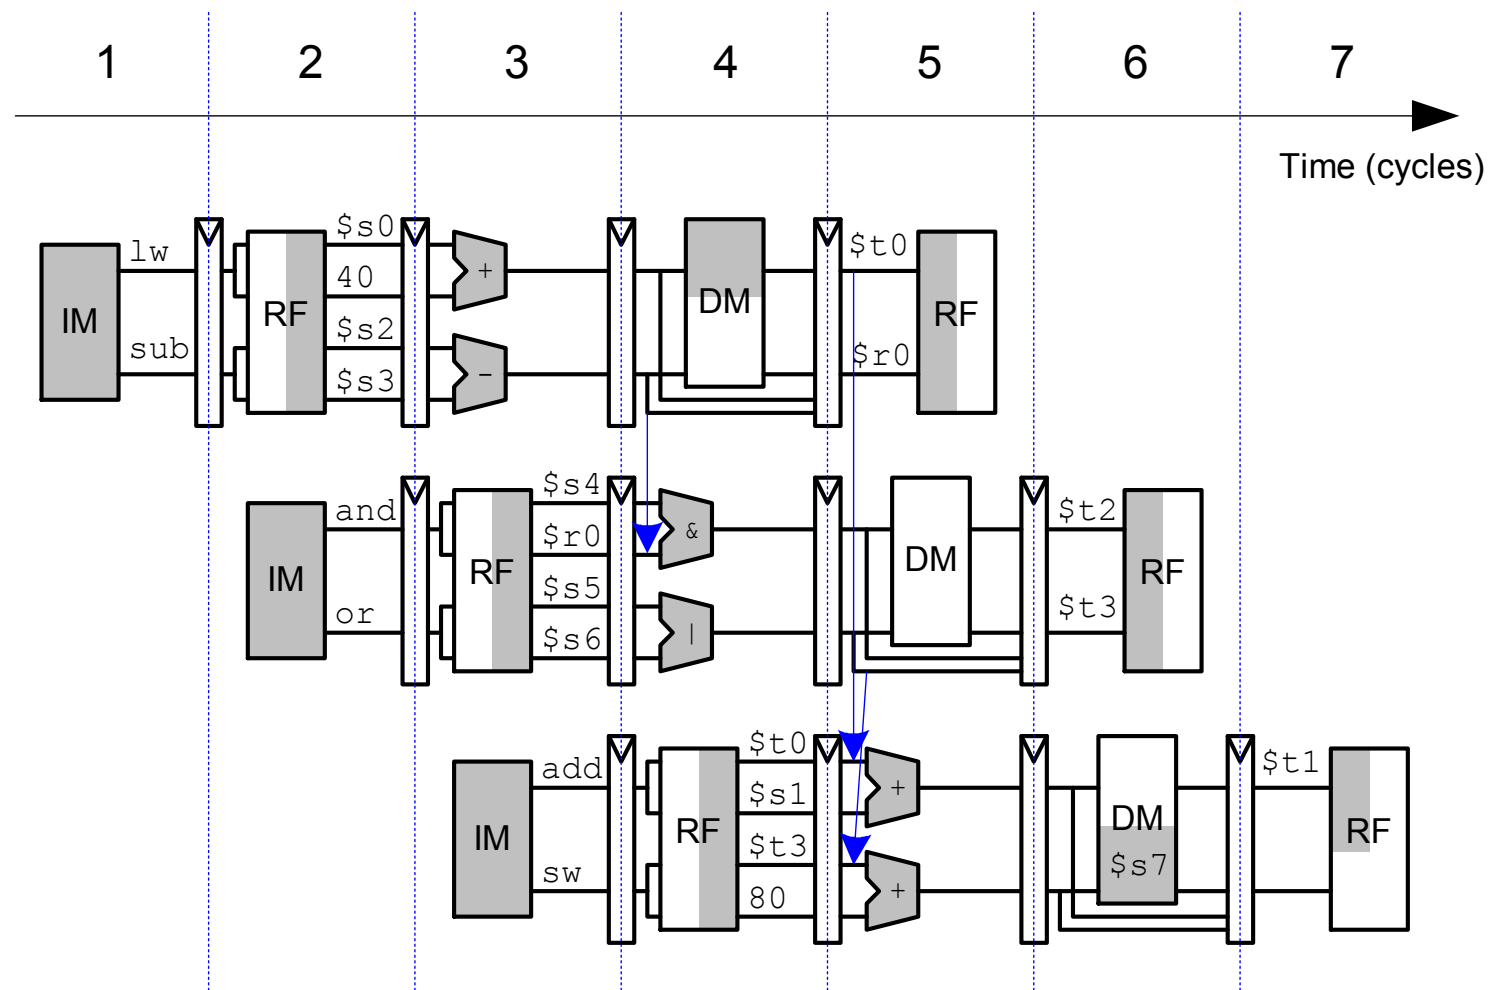
Register Renaming

```
lw    $t0, 40($s0)
add   $t1, $t0, $s1
sub   $t0, $s2, $s3
and   $t2, $s4, $t0
or    $t3, $s5, $s6
sw    $s7, 80($t3)
```

Ideal IPC: 2

Actual IPC: $6/3 = 2$

lw \$t0, 40(\$s0)
sub \$r0, \$s2, \$s3
2-cycle RAW and \$t2, \$s4, **RAW** \$r0
or \$t3, \$s5, \$s6
RAW add \$t1, \$t0, \$s1
sw \$s7, 80(\$t3)



Register Renaming

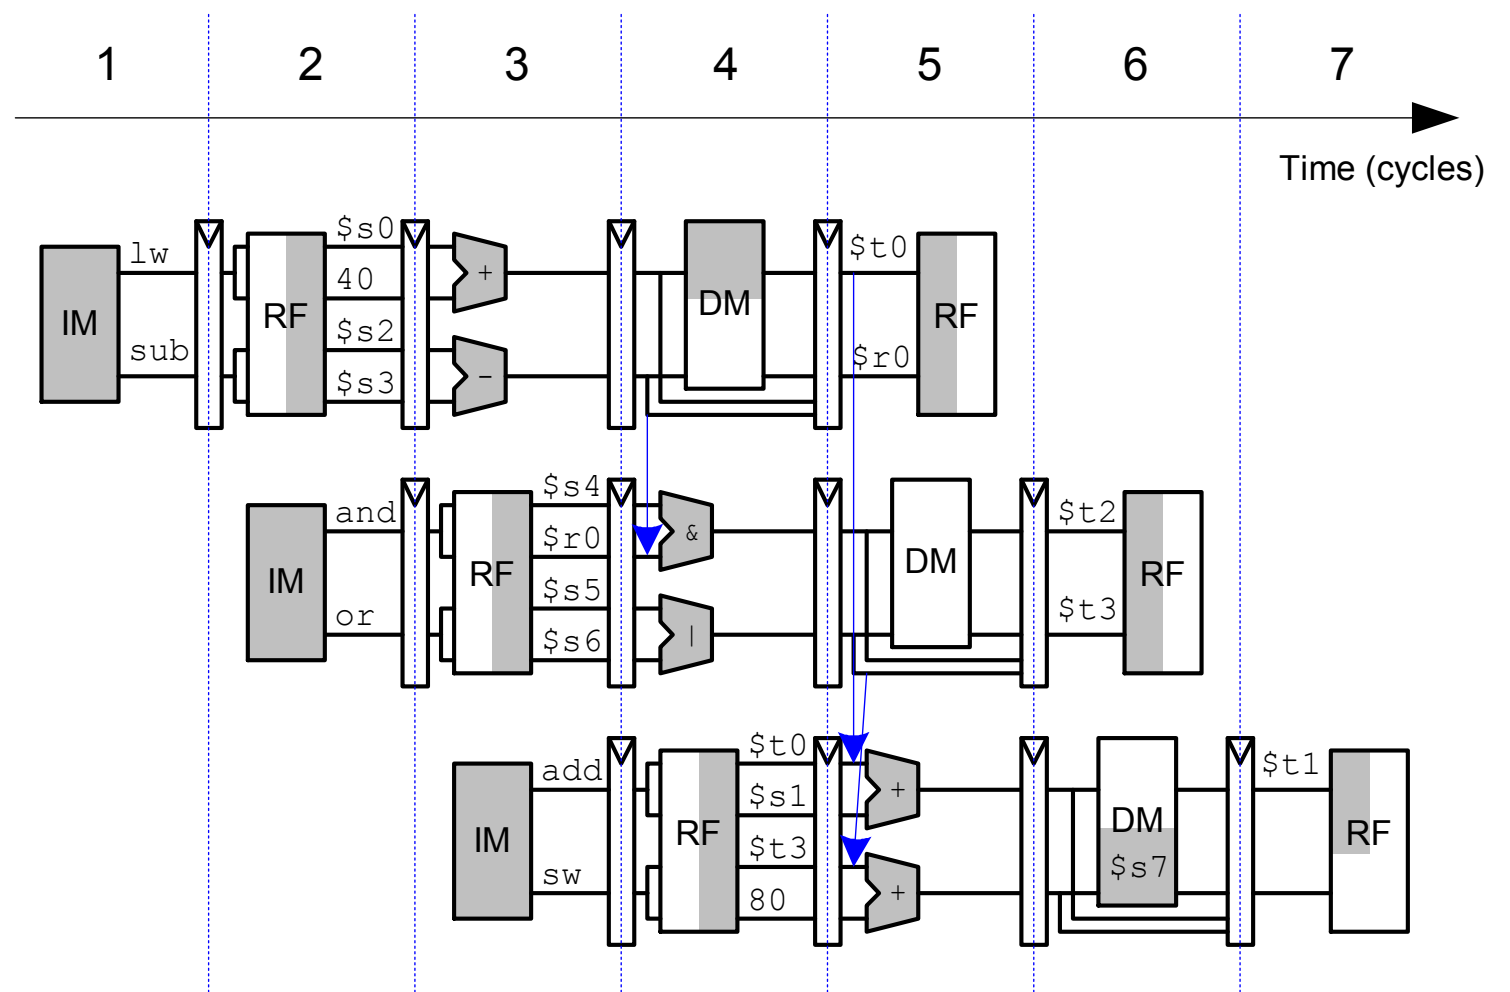
```
lw    $t0, 40($s0)
add   $t1, $t0, $s1
sub   $t0, $s2, $s3
and   $t2, $s4, $t0
or    $t3, $s5, $s6
sw    $s7, 80($t3)
```

r: renaming registers

Ideal IPC: 2

Actual IPC: $6/3 = 2$

lw \$t0, 40(\$s0)
sub \$r0, \$s2, \$s3
2-cycle RAW and \$t2, \$s4, **RAW** (\$r0)
or \$t3, \$s5, \$s6
RAW add \$t1, **RAW** (\$t0), \$s1
sw \$s7, 80(\$t3)

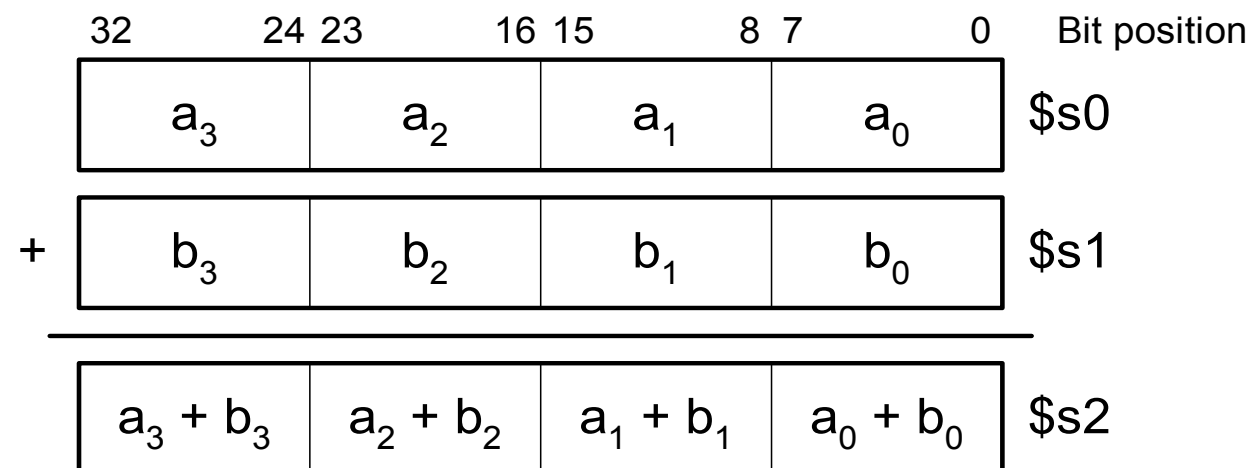


Some Advanced Concepts

SIMD

- ❖ Single Instruction Multiple Data (SIMD)
 - ❖ Single instruction acts on multiple pieces of data at once
 - ❖ Common application: graphics (8-bit RGB)
 - ❖ Perform short arithmetic operations (also called *packed arithmetic*)
- ❖ For example, add four 8-bit elements

`padd8 $s2, $s0, $s1`



Advanced Architecture Techniques

- ❖ **Multithreading**

- ❖ Wordprocessor: thread for typing, spell checking, printing
- ❖ More than one copy of the architectural state

- ❖ **Multiprocessors**

- ❖ Multiple processors (cores) on a single chip

Threading: Definitions

- ❖ **Process:** program running on a computer
 - ❖ Multiple processes can run at once: e.g., surfing Web, playing music, writing a paper
- ❖ **Thread:** part of a program
 - ❖ Each process has multiple threads: e.g., a word processor may have threads for typing, spell checking, printing

Threads in Conventional Processor

- ❖ One thread runs at once
- ❖ When one thread stalls (for example, waiting for memory):
 - ❖ Architectural state of that thread stored
 - ❖ Architectural state of waiting thread loaded into processor and it runs
 - ❖ Called **context switching**
- ❖ Appears to user like all threads running simultaneously

Multithreading

- ❖ Multiple copies of architectural state
- ❖ Multiple threads **active** at once:
 - ❖ When one thread stalls, another runs immediately
 - ❖ If one thread can't keep all execution units busy, another thread can use them
- ❖ Does not increase instruction-level parallelism (ILP) of single thread, but increases throughput

Intel calls this “hyperthreading”

Multiprocessors

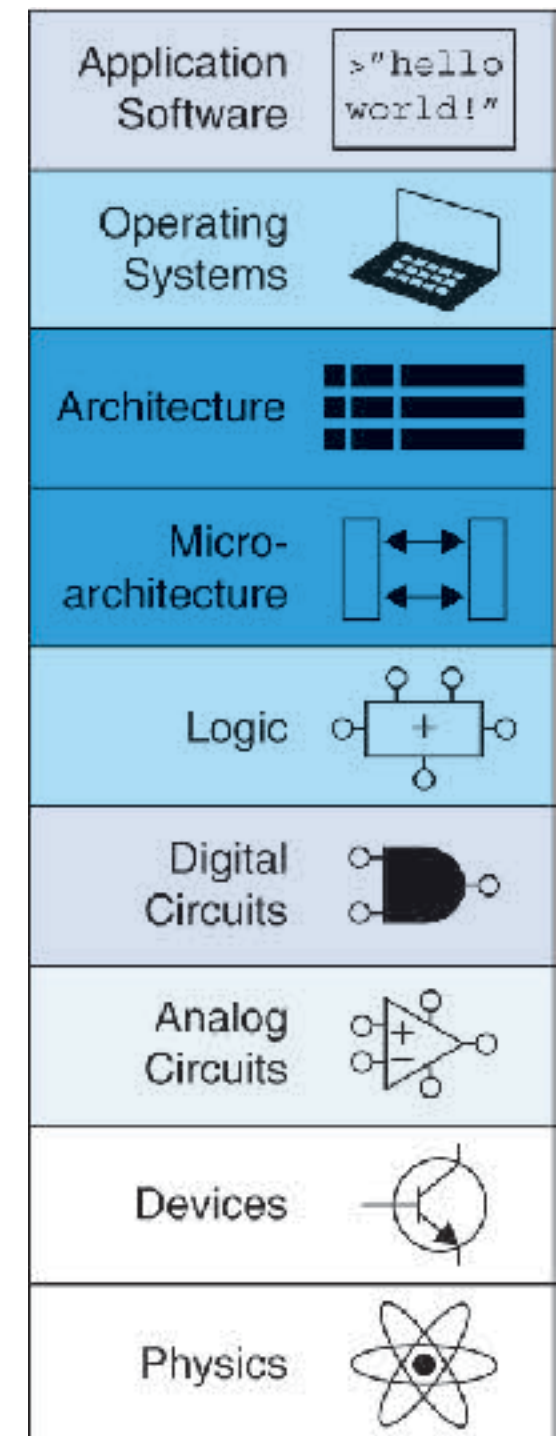
- ❖ Multiple processors (cores) with a method of communication between them
- ❖ Types:
 - ❖ **Homogeneous:** multiple cores with shared memory
 - ❖ **Heterogeneous:** separate cores for different tasks (for example, DSP and CPU in cell phone)
 - ❖ **Clusters:** each core has own memory system

Other Resources

- ❖ Patterson & Hennessy's: *Computer Architecture: A Quantitative Approach*
- ❖ Conferences:
 - ❖ www.cs.wisc.edu/~arch/
 - ❖ ISCA (International Symposium on Computer Architecture)
 - ❖ HPCA (International Symposium on High Performance Computer Architecture)

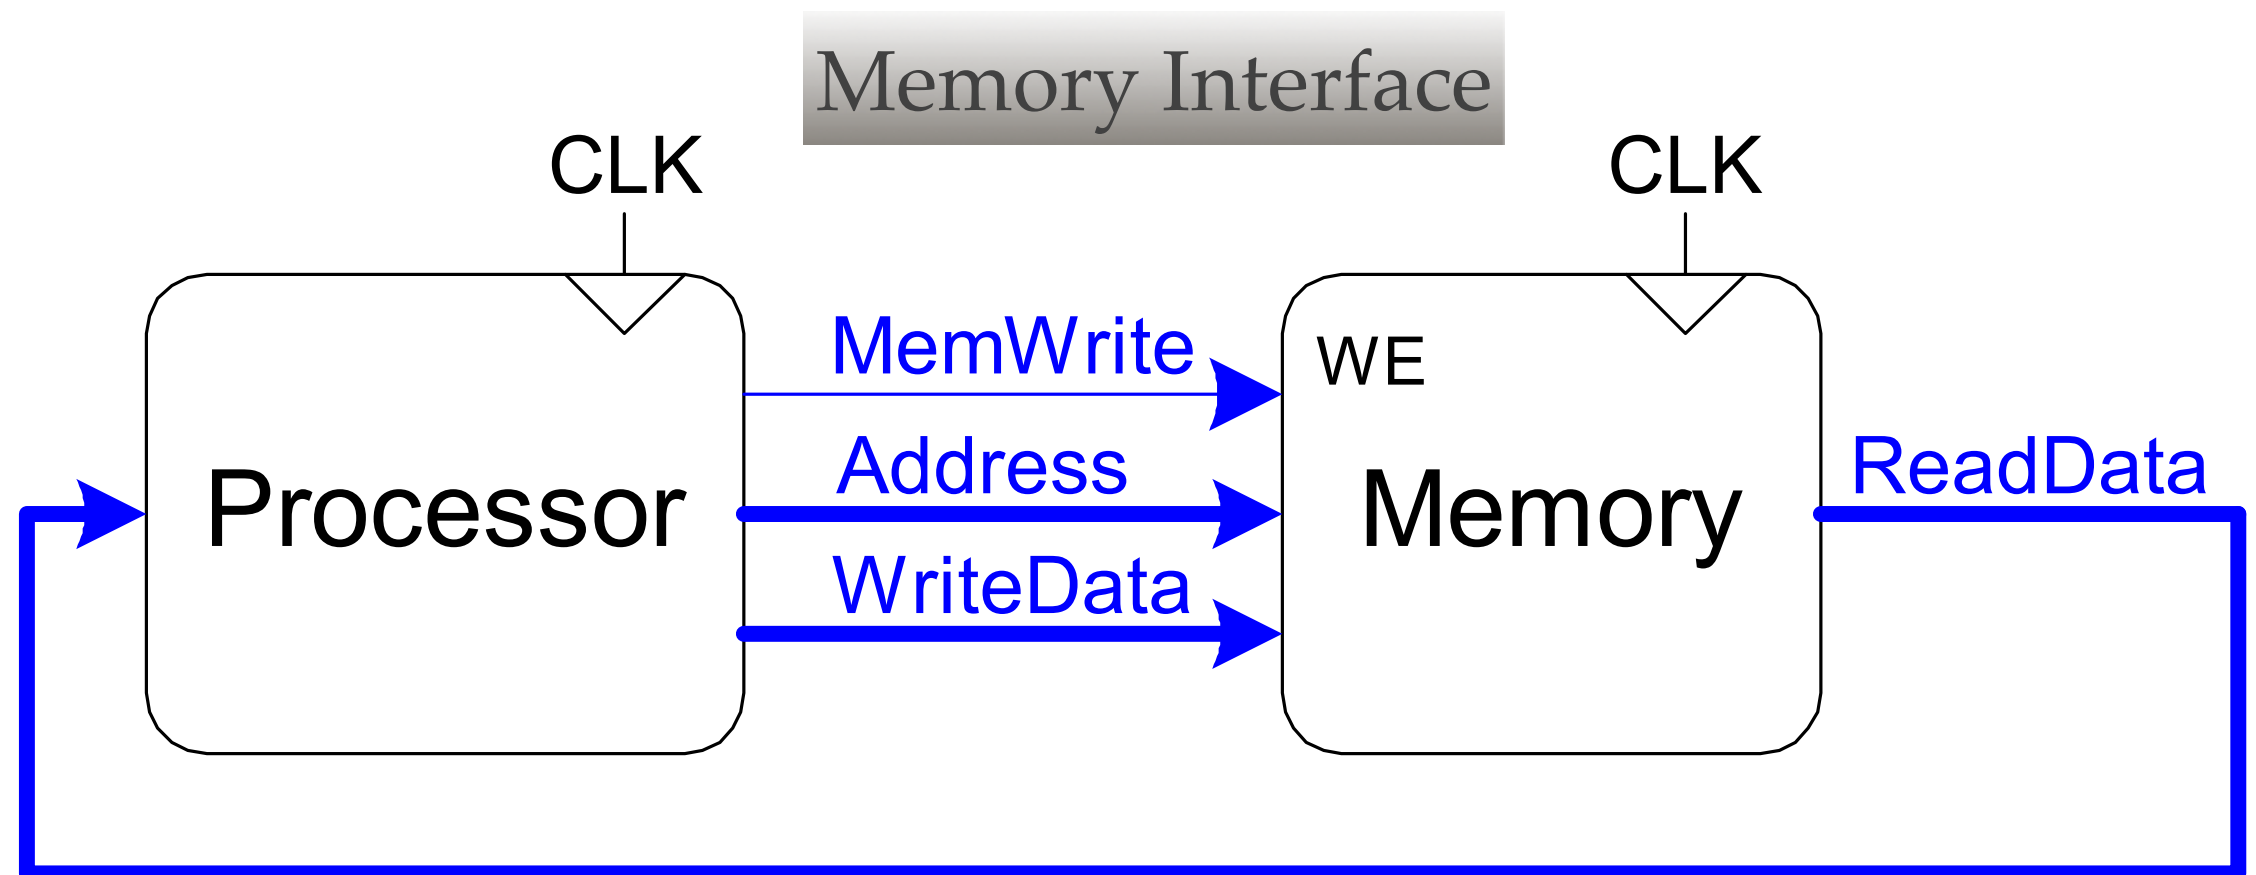
Other Related Topics

- ❖ Memory System Performance Analysis
- ❖ Caches
- ❖ Virtual Memory
- ❖ Memory-Mapped I/O




Memory as an Important Component

- ❖ Computer performance depends on:
 - ❖ Processor performance
 - ❖ Memory system performance



Processor-Memory Gap

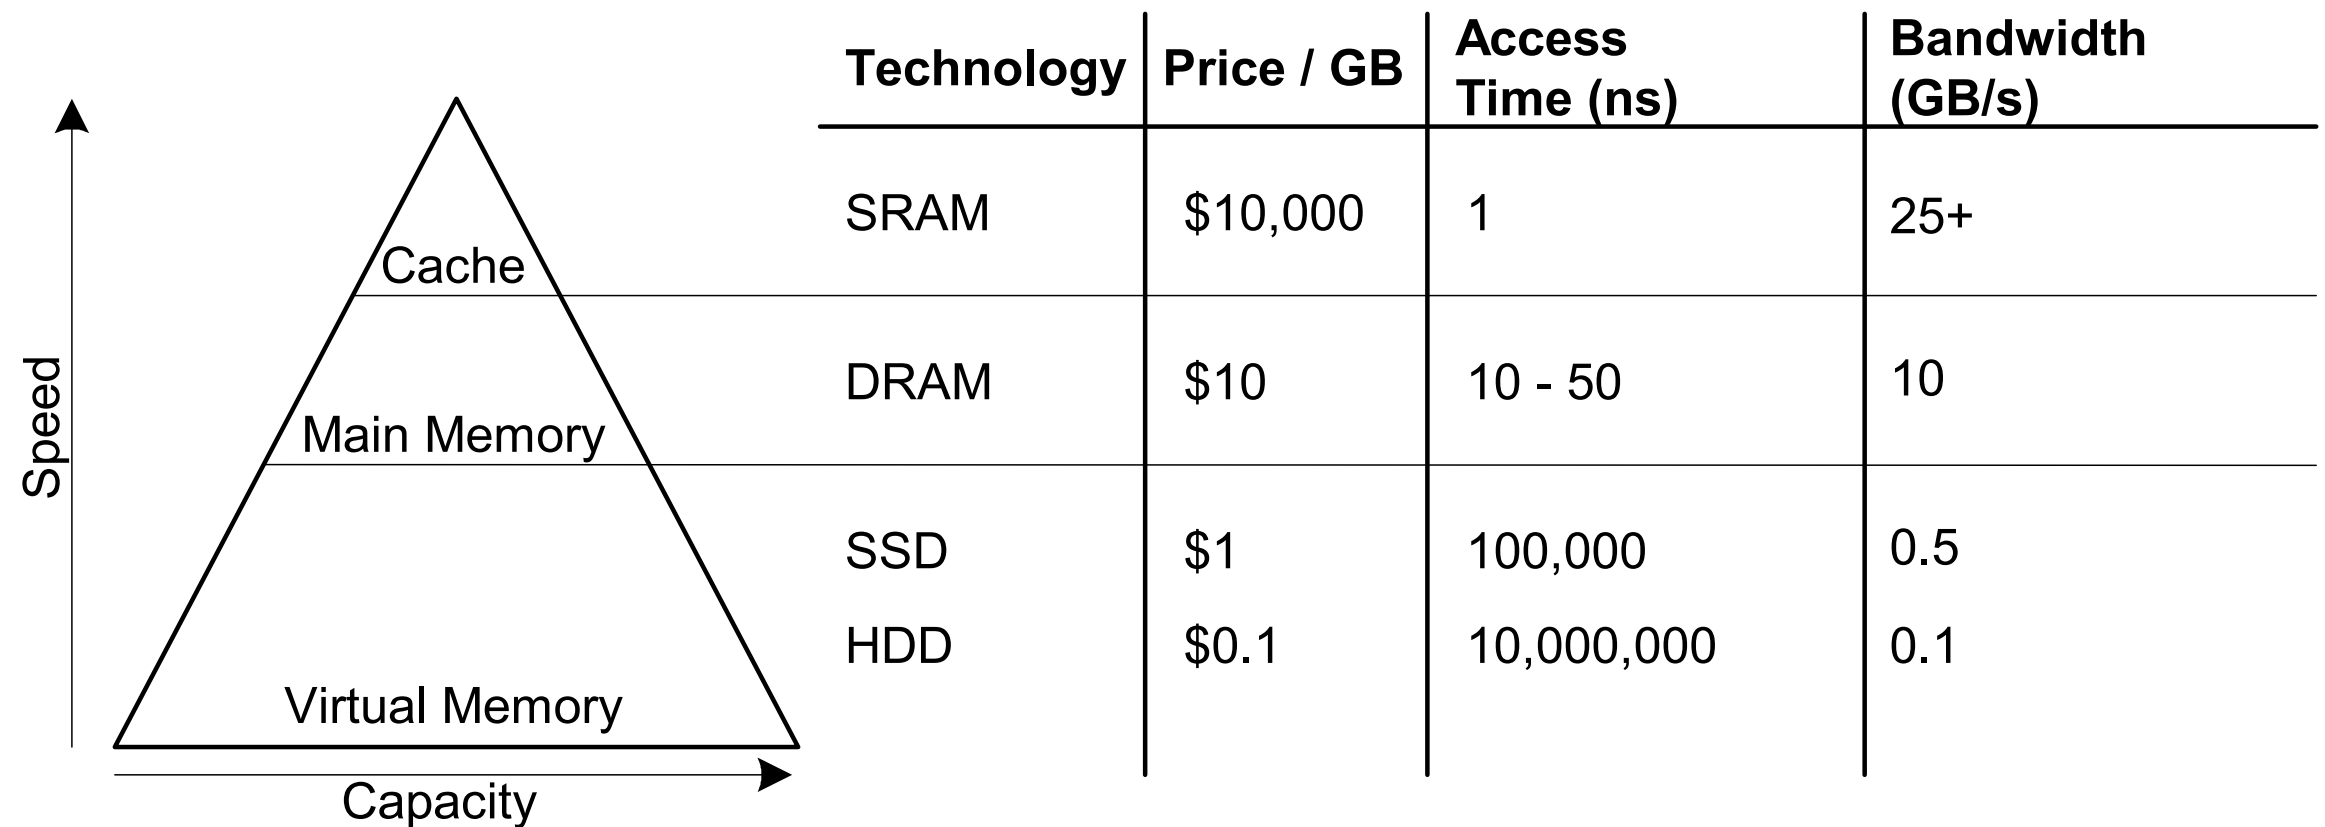
In prior chapters, assumed access memory in 1 clock cycle – but hasn't been true since the 1980's



Memory System Challenge

- ❖ Make memory system appear as fast as processor
 - ❖ Use hierarchy of memories
- ❖ Ideal memory:
 - ❖ Fast
 - ❖ Cheap (inexpensive)
 - ❖ Large (capacity)
- ❖ **But one can only has two!**

Memory Hierarchy



Locality

Exploit locality to make memory accesses fast

- ❖ **Temporal Locality:**

- ❖ Locality in time
- ❖ If data used recently, likely to use it again soon
- ❖ **How to exploit:** keep recently accessed data in higher levels of memory hierarchy

- ❖ **Spatial Locality:**

- ❖ Locality in space
- ❖ If data used recently, likely to use nearby data soon
- ❖ **How to exploit:** when access data, bring nearby data into higher levels of memory hierarchy too

Memory Performance

Hit: data found in that level of memory hierarchy

Miss: data not found (must go to next level)

$$\begin{aligned}\text{Hit Rate} &= \# \text{ hits} / \# \text{ memory accesses} \\ &= 1 - \text{Miss Rate}\end{aligned}$$

$$\begin{aligned}\text{Miss Rate} &= \# \text{ misses} / \# \text{ memory accesses} \\ &= 1 - \text{Hit Rate}\end{aligned}$$

Average memory access time (AMAT): average time for processor to access data

$$\text{AMAT} = t_{\text{cache}} + MR_{\text{cache}}[t_{MM} + MR_{MM}(t_{VM})]$$

Memory Performance Example 1

- ❖ A program has 2,000 loads and stores 1,250 of these data values in cache
- ❖ Rest supplied by other levels of memory hierarchy
- ❖ What are the hit and miss rates for the cache?
- ❖ **Hit Rate = $1250/2000 = 0.625$**
- ❖ **Miss Rate = $750/2000 = 0.375 = 1 - \text{Hit Rate}$**

Memory Performance Example 2

Suppose processor has 2 levels of hierarchy: cache and main memory

$$t_{\text{cache}} = 1 \text{ cycle}, t_{\text{MM}} = 100 \text{ cycles}$$

What is the AMAT of the program from Example 1?

Memory Performance Example 2

Suppose processor has 2 levels of hierarchy: cache and main memory

$$t_{\text{cache}} = 1 \text{ cycle}, t_{\text{MM}} = 100 \text{ cycles}$$

What is the AMAT of the program from Example 1?

$$\begin{aligned} \text{AMAT} &= t_{\text{cache}} + MR_{\text{cache}}(t_{\text{MM}}) \\ &= [1 + 0.375(100)] \text{ cycles} \\ &= \mathbf{38.5 \text{ cycles}} \end{aligned}$$

Cache

- ❖ Highest level in memory hierarchy
- ❖ Fast (typically ~ 1 cycle access time)
- ❖ Ideally supplies most data to processor
- ❖ Usually holds most recently accessed data

Cache Design Questions

- ❖ What data is held in the cache?
- ❖ How is data found?
- ❖ What data is replaced?

Focus on data loads, but stores follow same principles

What data is held in the cache?

- ❖ Ideally, cache anticipates needed data and puts it in cache
- ❖ But **impossible to predict future**
- ❖ Use past to predict future – temporal and spatial locality:
 - ❖ **Temporal locality:** copy newly accessed data into cache
 - ❖ **Spatial locality:** copy neighboring data into cache too

Cache Terminology

- ❖ **Capacity (C):**

- ❖ number of data bytes in cache

- ❖ **Block size (b):**

- ❖ bytes of data brought into cache at once

- ❖ **Number of blocks ($B = C/b$):**

- ❖ number of blocks in cache: $B = C / b$

- ❖ **Degree of associativity (N):**

- ❖ number of blocks in a set

- ❖ **Number of sets ($S = B/N$):**

- ❖ each memory address maps to exactly one cache set

How is data found?

- ❖ Cache organized into S sets
- ❖ Each memory address maps to exactly one set
- ❖ Caches categorized by # of blocks in a set:
 - ❖ **Direct mapped:** 1 block per set
 - ❖ **N -way set associative:** N blocks per set
 - ❖ **Fully associative:** all cache blocks in 1 set
- ❖ Examine each organization for a cache with:
 - ❖ Capacity ($C = 8$ words)
 - ❖ Block size ($b = 1$ word)
 - ❖ So, number of blocks ($B = 8$)

Example Cache Parameters

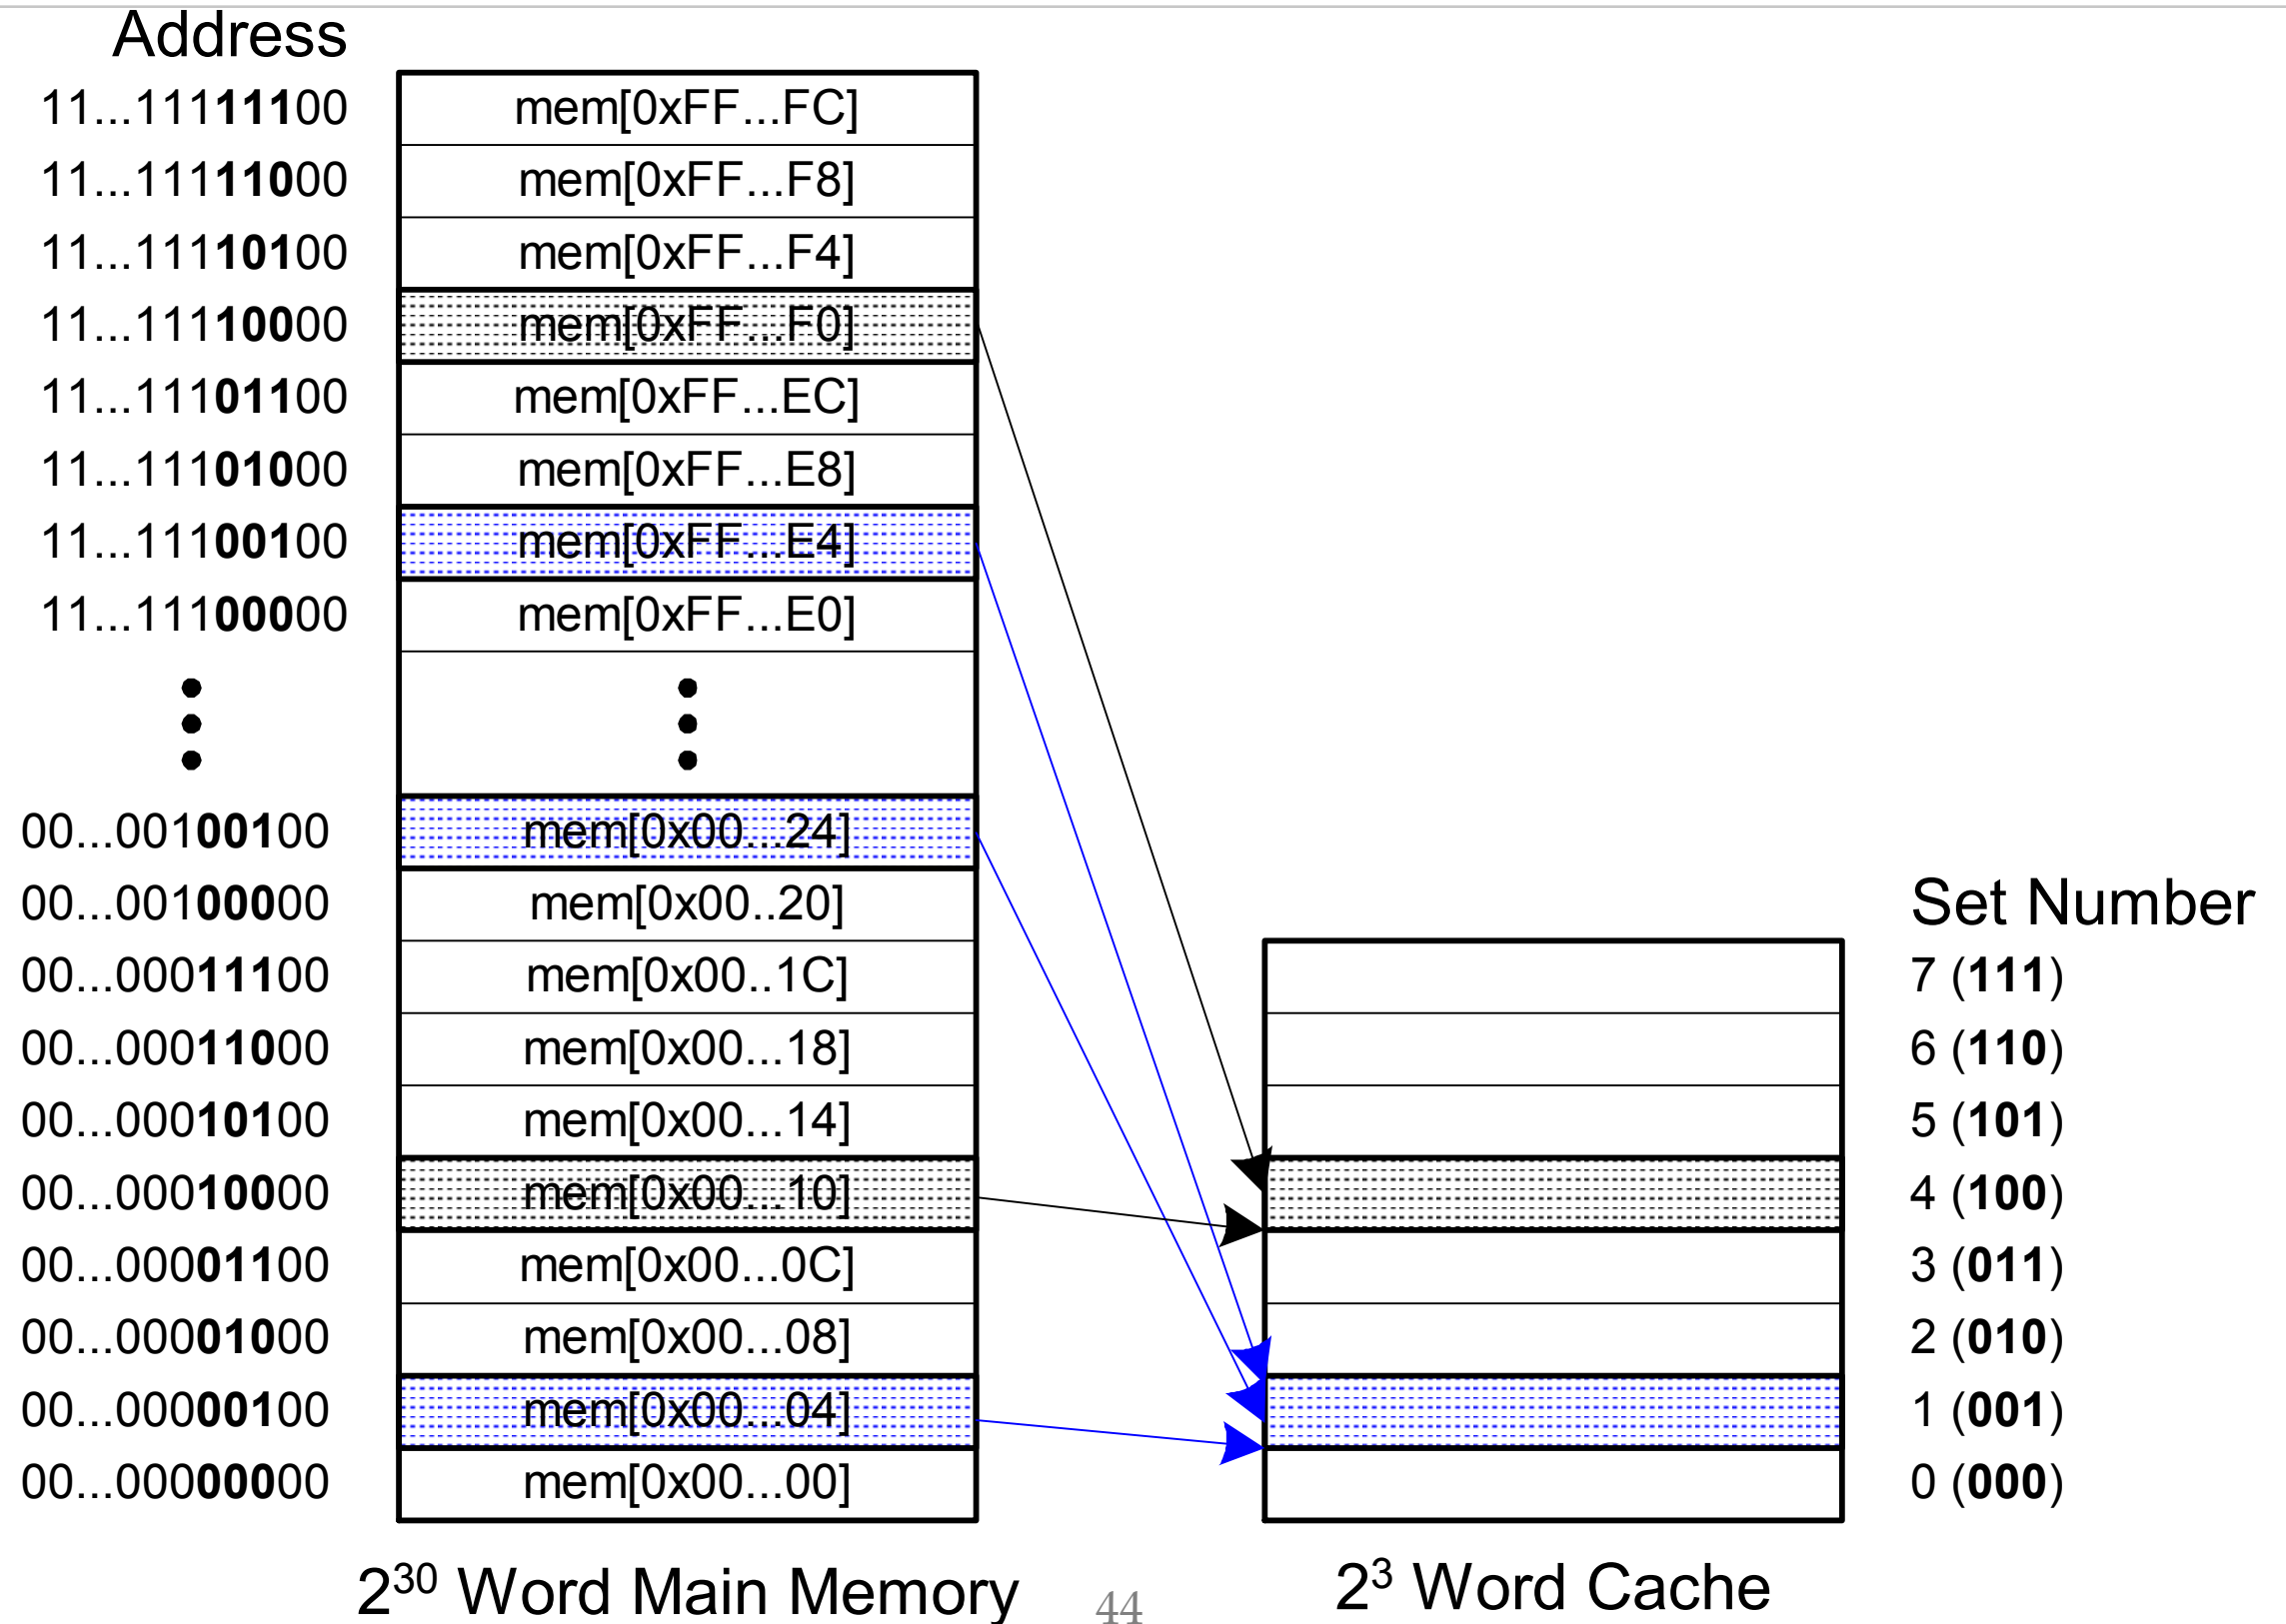
- ❖ $C = 8$ words (capacity)
- ❖ $b = 1$ word (block size)
- ❖ So, $B = 8$ (# of blocks)

Ridiculously small, but will illustrate organizations


Mapping

- ❖ Mapping: the relationship **between the address of data in main memory and the location of that data in the cache**
- ❖ Each memory address maps to exactly one set in the cache
 - ❖ Some of the address bits are used to determine which cache set contains the data


Direct Mapped Cache



Direct Mapped Cache Hardware



Direct Mapped Cache Performance



3

V Tag Data

0		
0		
0		
0		
1	00...00	mem[0x00...0C]
1	00...00	mem[0x00...08]
1	00...00	mem[0x00...04]
0		

Set 7 (111)

Set 6 (110)

Set 5 (101)

Set 4 (100)

Set 3 (011)

Set 2 (010)

Set 1 (001)


Set 0 (000)

MIPS assembly code

```
    addi $t0, $0, 5
loop: beq  $t0, $0, done
      lw   $t1, 0x4($0)
      lw   $t2, 0xC($0)
      lw   $t3, 0x8($0)
      addi $t0, $t0, -1
      j    loop
done:
```

Miss Rate = ?

Direct Mapped Cache Performance



3

V Tag Data

0		
0		
0		
0		
1	00...00	mem[0x00...0C]
1	00...00	mem[0x00...08]
1	00...00	mem[0x00...04]
0		

Set 7 (111)

Set 6 (110)

Set 5 (101)

Set 4 (100)

Set 3 (011)

Set 2 (010)

Set 1 (001)

Set 0 (000)

MIPS assembly code

```
    addi $t0, $0, 5
loop: beq  $t0, $0, done
      lw   $t1, 0x4($0)
      lw   $t2, 0xC($0)
      lw   $t3, 0x8($0)
      addi $t0, $t0, -1
      j    loop
done:
```

Miss Rate = 3/15
= 20%

Temporal Locality
Compulsory Misses

Direct Mapped Cache: Conflict

MIPS assembly code

```
    addi $t0, $0, 5
loop: beq  $t0, $0, done
      lw   $t1, 0x4($0)
      lw   $t2, 0x24($0)
      addi $t0, $t0, -1
      j    loop
done:
```

Memory Address	Tag	Set	Byte Offset
	00...01	001	00

3

V Tag Data

V	Tag	Data
0		
0		
0		
0		
0		
0		
1	00...00	mem[0x00...04] mem[0x00...24]
0		

Set 7 (111)
Set 6 (110)
Set 5 (101)
Set 4 (100)
Set 3 (011)
Set 2 (010)
Set 1 (001)
Set 0 (000)

Miss Rate = ?

Direct Mapped Cache: Conflict

MIPS assembly code

```
    addi $t0, $0, 5
loop: beq  $t0, $0, done
      lw   $t1, 0x4($0)
      lw   $t2, 0x24($0)
      addi $t0, $t0, -1
      j    loop
done:
```

Memory Address	Tag	Set	Byte Offset
	00...01	001	00

3

V Tag Data


0		
0		
0		
0		
0		
0		
1	00...00	mem[0x00...04] mem[0x00...24]
0		

Set 7 (111)
Set 6 (110)
Set 5 (101)
Set 4 (100)
Set 3 (011)
Set 2 (010)
Set 1 (001)
Set 0 (000)


Miss Rate = 10/10
= 100%

Conflict Misses

N-Way Set Associative Cache



N-Way Set Associative Cache



N-Way Set Associative Performance

MIPS assembly code

Miss Rate = ?

```
    addi $t0, $0, 5
loop: beq  $t0, $0, done
      lw   $t1, 0x4($0)
      lw   $t2, 0x24($0)
      addi $t0, $t0, -1
      j    loop
done:
```

Way 1			Way 0			
V	Tag	Data	V	Tag	Data	
0			0			Set 3
0			0			Set 2
0			0			Set 1
0			0			Set 0

N-Way Set Associative Performance

MIPS assembly code

```
    addi $t0, $0, 5
loop: beq  $t0, $0, done
      lw   $t1, 0x4($0)
      lw   $t2, 0x24($0)
      addi $t0, $t0, -1
      j    loop
done:
```

Miss Rate = 2/10
= 20%

Associativity reduces
conflict misses

Way 1			Way 0			
V	Tag	Data	V	Tag	Data	
0			0			Set 3
0			0			Set 2
1	00...10	mem[0x00...24]	1	00...00	mem[0x00...04]	Set 1
0			0			Set 0

Fully Associative Cache

V	Tag	Data	V	Tag	Data	V	Tag	Data	V	Tag	Data	V	Tag	Data	V	Tag	Data	V	Tag	Data

Reduces conflict misses

Expensive to build